

surfacehop_jax: User Guide

Version 1.1.0 -- Differentiable FSSH in JAX

Prateek Goel

May 2026

Contents

surfacehop_jax — User Guide	3
Overview	3
Where to start	3
Table of contents	3
Conventions used in this guide	4
Other resources	4
Chapter 1: Introduction	4
What surface hopping is	4
What this package provides	5
When to use surfacehop_jax	5
What’s in the package	6
Chapter 2: Installation	6
Requirements	6
From a clone	6
GPU install	7
Testing the install	7
Common install issues	7
Next	8
Chapter 3: Theory — the FSSH algorithm	8
3.1 Diabatic vs adiabatic representations	8
3.2 The two equations of motion	9
3.3 The fewest-switches hop probability	10
3.4 Momentum rescaling on a hop	10
3.5 The eigenvector phase problem	11
3.6 Internal consistency	12
3.7 What FSSH gets wrong, and what to do about it	12
Next	12
Chapter 4: Quickstart	12
4.1 A single Tully-1 trajectory	13
4.2 A pyrazine ensemble	13
4.3 A differentiable parameter fit	14
Next	16
Chapter 5: Models — built-in Hamiltonians and how to write your own	16

5.1 The model interface	16
5.2 Tully 1990 models	16
5.3 Linear vibronic coupling	17
5.4 The pyrazine 4-mode benchmark	19
5.5 Writing a custom Hamiltonian	19
5.6 ML potentials	20
Next	21
Chapter 6: API reference	21
6.1 Module structure	21
6.2 Constants	21
6.3 Models (surfacehop_jax.models)	21
6.4 PES (surfacehop_jax.pes)	22
6.5 Dynamics (surfacehop_jax.dynamics)	23
6.6 Wigner sampling (surfacehop_jax.wigner)	24
6.7 Decoherence (surfacehop_jax.decoherence)	25
Next	25
Chapter 7: Wigner sampling of initial conditions	25
7.1 The ground-state Wigner function	25
7.2 The API	26
7.3 Why this matters: the mass-factor bug	26
7.4 Sampling in dimensionless coordinates	27
7.5 Wigner sampling for an LVC: full pattern	27
7.6 Sampling around a non-equilibrium geometry	28
7.7 Reproducibility	28
Next	28
Chapter 8: Decoherence corrections	28
8.1 What “over-coherence” means in FSSH	28
8.2 The Granucci-Persico EDC: physics	29
8.3 Using EDC	30
8.4 Tuning α (the kinetic-energy floor)	30
8.5 When EDC is and isn’t necessary	31
8.6 Plug-in custom decoherence	31
Next	32
Chapter 9: Differentiable workflows	32
9.1 The differentiable propagator: what flows through <code>jax.grad</code>	32
9.2 The “pathwise” gradient: smooth observables vs hop counts	33
9.3 Worked recipe: fit one parameter	33
9.4 Multi-parameter gradients in one backward pass	34
9.5 Validation against finite difference	35
9.6 Common pitfalls	35
9.7 Beyond LVC fits: ideas	36
Next	36
Chapter 10: Performance — JIT, vmap, GPU	37
10.1 The cost model	37
10.2 Why JIT helps	37
10.3 Why vmap helps even more	38
10.4 Compile time and trace caching	38
10.5 GPU	39
10.6 Memory	39
10.7 When FSSH gets slow	39

10.8 Profiling	40
Next	40
Chapter 11: Troubleshooting	40
11.1 Setup and import issues	40
11.2 Forward-pass problems	41
11.3 Ensemble problems	42
11.4 Autodiff problems	42
11.5 Numerical / accuracy issues	43
11.6 Getting help	44
About this guide	44

surfacehop_jax — User Guide

Version 1.1.0 • Differentiable fewest-switches surface hopping in JAX

Overview

This directory is the long-form documentation for surfacehop_jax. The top-level [README.md](#) is a quick advert; the JOSS paper [joss/paper.md](#) is the publication. This guide is for **users who have cloned the repository and want to actually do something** — run a benchmark, build their own LVC model, fit a parameter against an experimental observable, or understand why a particular FSSH result looks the way it does.

Where to start

If you want to ...	Read these in order
Understand what surface hopping is and why this package exists	01 — Introduction → 03 — Theory
Just get a calculation running	02 — Installation → 04 — Quickstart
Set up a real photochemistry model	04 — Quickstart → 05 — Models → 07 — Wigner sampling
Use the decoherence correction	08 — Decoherence
Fit LVC parameters to an observable with <code>jax.grad</code>	09 — Differentiable workflows and the notebook differentiable_dynamics.ipynb
Run large ensembles on a GPU	10 — Performance
Look up a function	06 — API reference
Diagnose a crash or a wrong-looking number	11 — Troubleshooting

Table of contents

- [1. Introduction — what surface hopping is and why this package exists](#)
- [2. Installation and testing](#)
- [3. Theory — the FSSH algorithm in detail](#)
- [4. Quickstart — three working examples](#)
- [5. Models — built-in and custom Hamiltonians](#)
- [6. API reference](#)
- [7. Wigner sampling of initial conditions](#)
- [8. Decoherence corrections](#)

9. [Differentiable workflows — `jax.grad` through trajectories](#)
10. [Performance — JIT, vmap, GPU](#)
11. [Troubleshooting](#)

Conventions used in this guide

- **Atomic units everywhere.** Energies in Hartree, masses in electron masses, distances in Bohr, time in atomic units of time (1 a.u. = 0.0242 fs). The constants module exports the conversions if you need them: `surfacehop_jax.constants.HARTREE_TO_EV`, `AU_OF_TIME_TO_FS`, etc.
- **All array shapes are spelled out.** `(nel,)` means an array indexed by electronic state, `(ndim,)` by nuclear coordinate, `(n_traj,)` by trajectory in an ensemble, and so on.
- **Code examples are runnable as written**, given a working install and the imports shown at the top of each section.
- **Mathematical equations** use LaTeX. GitHub now renders inline `$...$` and display `$$...$$` math directly in markdown.

Other resources

- The Jupyter notebook [notebooks/differentiable_dynamics.ipynb](#) is the canonical end-to-end demonstration of gradient-based parameter fitting.
- The example scripts in [examples/](#) reproduce the figures in the JOSS paper.
- The test suite under [tests/](#) is also useful as documentation — every public function has at least one short, readable test that exercises it.

Chapter 1: Introduction

What surface hopping is

Most of computational chemistry rests on the **Born-Oppenheimer approximation**: nuclei are heavy and slow, electrons are light and fast, so we solve for the electrons at fixed nuclear positions to get a single potential-energy surface $V(\mathbf{R})$ and then propagate the nuclei on that surface. This works beautifully for ground-state chemistry. It fails — sometimes spectacularly — when two electronic states come close in energy and the nuclear motion can switch between them. Photochemistry is full of these situations:

- A molecule absorbs a UV photon, lands on an excited surface, slides toward a region where that surface kisses a lower one (a *conical intersection*), and decays radiationlessly to the ground state in tens of femtoseconds.
- A solar-cell heterojunction transfers a photogenerated electron from one chromophore to its neighbour by hopping between adjacent diabatic states whenever the gap closes.
- A photoisomerisation (vision in the retina, photoswitches in molecular machines) is a non-adiabatic transition along a torsional coordinate.

For these problems we need **nonadiabatic** molecular dynamics: nuclei that can change electronic state during their motion. The exact treatment is the full time-dependent Schrödinger equation in the nuclear-electronic space, which beyond a handful of degrees of freedom is intractable. Tully’s *fewest-switches surface hopping* (FSSH, introduced in [Tully 1990](#)) is the most widely used classical-trajectory approximation: each trajectory follows Newton’s equations on **one** adiabatic surface at any instant, but is allowed to switch (“hop”) to a neighbouring surface stochastically, with a probability constructed so that the swarm of trajectories reproduces the populations given by the time-dependent Schrödinger equation.

The “fewest-switches” criterion is what makes FSSH efficient: the hop probability is chosen to be the *minimum* number of hops consistent with maintaining the correct ensemble populations. Naive “always-hop-when-coupled” schemes overcount transitions; pure mean-field methods (Ehrenfest) underweight the influence of strong but brief couplings. FSSH threads the needle.

What this package provides

`surfacehop_jax` is a clean re-implementation of FSSH in the JAX numerical framework. Three things follow from that choice:

1. **JIT compilation.** The full propagator — velocity-Verlet for nuclei, exact matrix exponential for the electronic TDSE, hopping decision, momentum rescaling on hop, frustrated-hop reflection — compiles to a single XLA program. There is no Python overhead per step.
2. **Vectorisation over an ensemble.** An ensemble of N trajectories is a single `jax.vmap` of the per-trajectory propagator. Once compiled it runs as one batched program: zero per-trajectory Python overhead, and transparent GPU execution for “free” 10k-trajectory ensembles on consumer hardware.
3. **End-to-end differentiability.** Every line of the propagator is a pure JAX function, so `jax.grad` can reverse-mode-differentiate any smooth observable of the final state with respect to any input parameter — a coupling strength, an excitation energy, a frequency, a Wigner-distribution width, the weights of an ML-trained potential. This is what unlocks the *inverse problem*: fit a model Hamiltonian to an experimental observable using exactly the gradient machinery that trains neural networks.

The package also ships convenience models that turn what is normally a multi-week setup task into one-liners:

- The three `TullyModel{1,2,3}` benchmarks from the 1990 paper. Useful for sanity checks, teaching, and validating any new feature.
- A general `LinearVibronicCoupling` class for the Köppel-Domcke-Cederbaum diabatic Hamiltonian. You supply vertical excitation energies, mode frequencies, and the κ, λ tensors; it returns a JAX-traceable Hamiltonian.
- A `pyrazine_4mode()` factory that returns the canonical 4-mode pyrazine S_1/S_2 benchmark used throughout the nonadiabatic-dynamics literature.

When to use `surfacehop_jax`

This is a *model Hamiltonian* code, not an on-the-fly ab initio one. You should reach for it when:

- you have an analytic or pre-fit Hamiltonian (LVC, Marcus-like bilinear, lattice models, ML-PES) and want fast ensemble dynamics on it;
- you want gradient-based fitting of model parameters against an experimental observable;
- you want to embed nonadiabatic dynamics inside a larger differentiable pipeline (e.g. inverse design of a chromophore for a target decay rate, or training an ML potential against quantum reference trajectories);
- you want a small, readable, hackable FSSH propagator (~400 lines) rather than a hundred-thousand-line legacy package.

You should reach for SHARC, Newton-X, JADE, or PYXAID instead when:

- you need on-the-fly couplings from a quantum-chemistry program (CASSCF, MS-CASPT2, ADC(2), LR-TDDFT);

- you need spin-orbit coupling, intersystem crossing, or relativistic effects;
- you need atom-feature outputs (geometries, dipoles, photoelectron spectra computed per step) for direct comparison to experiment.

In practice, model-Hamiltonian dynamics and on-the-fly dynamics are complementary: an LVC model fit to a few hundred ab initio points and then propagated by `surfacehop_jax` gives you a converged ensemble in seconds; an on-the-fly run gives you spectroscopic observables you could not parameterise into a model. The companion package `nma_jax` bridges the two: it produces the normal modes you need to define an LVC Hamiltonian in the first place.

What's in the package

```

surfacehop_jax/
├── surfacehop_jax/           # the package itself, ~1000 lines
│   ├── __init__.py          # public API
│   ├── constants.py         # HBAR, conversions
│   ├── models.py            # TullyModel{1,2,3}, LinearVibronicCoupling, pyrazine_4mode
│   ├── pes.py               # diabatic → adiabatic transform (energies, gradients, NACs)
│   ├── dynamics.py          # TrajectoryState, step, simulate, run_ensemble
│   ├── wigner.py            # H0 ground-state Wigner sampling
│   └── decoherence.py       # Granucci–Persico EDC corrections
├── tests/                   # 67 tests across all modules
├── examples/                # runnable example scripts
├── notebooks/               # Jupyter notebooks (parameter fitting)
├── docs/                    # this guide
└── joss/                     # JOSS paper sources

```

Onward to [installation](#).

Chapter 2: Installation

Requirements

- Python ≥ 3.10 (uses dataclasses, the new typing union syntax, and Self-style annotations).
- JAX ≥ 0.4 . `jaxlib` is bundled with `jax` for CPU; for GPU follow the official JAX install instructions (link below).
- `numpy`, `optax` (for the parameter-fit notebook), `matplotlib` (examples and the notebook).
- `pytest` for the test suite.

`surfacehop_jax` enables JAX's float64 mode on import (see `__init__.py`). Eigendecompositions of nearly-degenerate Hamiltonians and long propagations both lose precision in float32, so this is deliberate. If you mix `surfacehop_jax` with another package that assumes float32, you'll need to be careful about where each library promotes/demotes dtypes.

From a clone

The typical install for development or notebook use:

```

git clone https://github.com/mowgliamu/surfacehop_jax.git
cd surfacehop_jax
pip install -e .[test]

```

The editable (-e) install lets you patch the source and re-run tests without reinstalling. The [test] extras pull in pytest and matplotlib.

GPU install

JAX maintains the canonical GPU install instructions at <https://jax.readthedocs.io/en/latest/installation.html>. The short version on a Linux box with a recent NVIDIA driver:

```
pip install --upgrade "jax[cuda12]"
```

Then install surfacehop_jax on top:

```
pip install -e .
```

surfacehop_jax itself contains no GPU-specific code. Every function in the package is vmap- and jit-compatible; whether it runs on CPU or GPU is determined entirely by which jaxlib you have installed and which device the inputs live on. The same example scripts that run in a few seconds on CPU run an ensemble of thousands of trajectories in the same wall time on a single consumer GPU.

To confirm the GPU is being seen:

```
import jax
print(jax.devices())
# CPU-only: [CpuDevice(id=0)]
# GPU:      [CudaDevice(id=0)] (or similar)
```

If jax.devices() returns CpuDevice despite having installed the CUDA wheel, the most common culprit is a mismatched CUDA-toolkit / driver version. JAX's install page has a CUDA-version compatibility table.

Testing the install

```
pytest tests/           # fast subset, runs in ~4 minutes on a laptop CPU
pytest tests/ -m slow    # the Tully Model 1 momentum-scan benchmark; ~30 minutes on CPU
```

The default pytest tests/ skips tests marked slow. The slow group is the full Tully-1 momentum scan: 25 momenta × 500 trajectories each, with a tolerance check against a digitised reference curve. It is what we *publish* in the JOSS paper but is not required for everyday use.

If everything passes (61 fast tests, 6 slow tests = 67 total at v1.1) you're done.

Common install issues

jaxlib not found. Usually means pip resolved a jax version newer than the jaxlib it can find on your platform. Pinning both versions usually fixes it; see the JAX install matrix.

jax.config.update("jax_enable_x64", True) fails or is ignored. If you imported JAX before surfacehop_jax, JAX has already cached its default dtype as float32. The package enables x64 at import time, which is the right time — but if some upstream module disables it later, our eigendecompositions may give you garbage. Put import surfacehop_jax as early as practical, ideally before any other JAX-using import in your script.

ModuleNotFoundError: optax. The differentiable-fitting notebook uses optax for Adam. The minimal install does not pull it in automatically; pip install optax is enough.

On macOS, pip install -e . complains about setuptools. Newer pip defaults to PEP 517 builds. Either upgrade pip/setuptools or use pip install --use-pep517 -e . explicitly.

Next

If you want the conceptual picture before running anything, continue to [03 — Theory](#). If you’d rather see code first, jump to [04 — Quickstart](#).

Chapter 3: Theory — the FSSH algorithm

This chapter explains what every line of `surfacehop_jax.dynamics.step` actually does. If you just want to run a calculation, skip to [Quickstart](#) and come back later; nothing in the package’s public API requires reading this chapter. But if a number looks wrong, or you want to extend the propagator, this is where to start.

3.1 Diabatic vs adiabatic representations

For a system with n_{el} electronic states and a nuclear coordinate vector $\mathbf{Q}$ of dimension n_{dim} , the electronic Hamiltonian at a fixed nuclear geometry is an $n_{\text{el}} \times n_{\text{el}}$ matrix $\mathbf{H}(\mathbf{Q})$. You can pick the basis you write it in:

- The **diabatic** basis is whatever convenient basis you set up at a reference geometry — typically the eigenstates of $\mathbf{H}$ there. The diabatic Hamiltonian is then *smooth in $\mathbf{Q}$* away from the reference, but generally **not diagonal**: at other geometries the diabatic basis is no longer an eigenbasis of $\mathbf{H}$.
- The **adiabatic** basis is, at every $\mathbf{Q}$, the eigenbasis of $\mathbf{H}(\mathbf{Q})$. In this basis $\mathbf{H}$ is diagonal at every point — the diagonal entries are the *adiabatic surfaces* $E_i(\mathbf{Q})$ — but the basis itself rotates with $\mathbf{Q}$, so the time derivative picks up coupling terms $\mathbf{d}_{ij}(\mathbf{Q}) = \langle \psi_i(\mathbf{Q}) | \nabla_{\mathbf{Q}} \psi_j(\mathbf{Q}) \rangle$, the **non-adiabatic coupling vectors** (NACs).

`surfacehop_jax` takes the diabatic Hamiltonian as user input (it’s the natural form for an LVC or for an ML-fit potential) and computes adiabatic energies, gradients, and NACs internally via autodiff. The trajectory itself lives on adiabatic surfaces — this is what “surface” in “surface hopping” refers to.

How `adiabatic_quantities` works

Given any callable $H(\mathbf{x}) \rightarrow (n_{\text{el}}, n_{\text{el}})$ Hermitian matrix, `pes.adiabatic_quantities` returns the energies, gradients, NACs, and eigenvectors at a single point $\mathbf{x}$:

```
H = diab_h_fn(x)                # (nel, nel)
energies, eigvecs = jnp.linalg.eigh(H)  # (nel,), (nel, nel)
dH = jax.jacrev(diab_h_fn)(x)        # (nel, nel, ndim)
T = jnp.einsum("ai,abk,bj->ijk",
               eigvecs, dH, eigvecs)   # (nel, nel, ndim)
```

$T[i, j, k]$ is $\langle \psi_i | \partial H / \partial x_k | \psi_j \rangle$, the matrix element of the nuclear gradient of the (diabatic) Hamiltonian sandwiched between adiabatic eigenstates. By Hellmann–Feynman:

$$\frac{\partial E_i}{\partial x_k} = T_{iik},$$
$$\mathbf{d}_{ij}^{(k)} = \frac{T_{ijk}}{E_j - E_i} \quad (i \neq j).$$

That’s it — one eigh, one jacrev, and one einsum, and we have everything the FSSH propagator needs. The $1/(E_j - E_i)$ denominator is handled with the “double-where” pattern to avoid NaN gradients on the diagonal.

3.2 The two equations of motion

A single FSSH trajectory carries two pieces of state besides position and velocity:

1. **Active state** $s \in \{0, 1, \dots, n_{el}-1\}$: the adiabatic surface the nuclei are currently moving on. Integer.
2. **Electronic coefficients** $\mathbf{c}(t) \in \mathbb{C}^{n_{el}}$: the complex amplitudes of an electronic wavefunction $\Psi(t) = \sum_i c_i(t) \psi_i(\mathbf{Q}(t))$ propagated along the trajectory. Normalised: $\sum_i |c_i|^2 = 1$.

Nuclear EOM

Newton’s equations on the active surface:

$$m_k \ddot{Q}_k = -\frac{\partial E_s}{\partial Q_k}.$$

We integrate this with velocity Verlet, which is symplectic and so conserves energy on average between hops (drift is purely from the discrete hops themselves; see momentum rescaling below).

Electronic EOM

The amplitudes evolve under the time-dependent Schrödinger equation in the adiabatic basis, with the nuclear velocity coupling adiabatic states through their NACs:

$$i\hbar \dot{c}_i = E_i c_i - i\hbar \sum_j (\mathbf{v} \cdot \mathbf{d}_{ij}) c_j.$$

Define the generator

$$G_{ij} = -\frac{iE_i}{\hbar} \delta_{ij} - \mathbf{v} \cdot \mathbf{d}_{ij}.$$

Then $\dot{\mathbf{c}} = \mathbf{G}\mathbf{c}$, and we propagate over a nuclear step Δt by exact matrix exponential:

$$\mathbf{c}(t + \Delta t) = e^{\mathbf{G}\Delta t} \mathbf{c}(t).$$

This is `jax.scipy.linalg.expm(G * dt) @ coeffs` in the code. It is exact for constant $\mathbf{G}$ over the step (we hold the NACs and velocity fixed across $[t, t + \Delta t]$). Higher-order integrators that interpolate $\mathbf{G}(t)$ across the step would be a straightforward extension; the matrix-exponential single-step scheme is robust and sufficient for time steps $\Delta t \lesssim 1$ a.u. ≈ 0.024 fs in typical photochemistry problems.

Why matrix exponential instead of solve_ivp? The naive choice would be SciPy’s `solve_ivp` with RK45 or LSODA. Both work, but complex-valued ODE support in SciPy is poor (you have to split into real/imaginary parts), and neither is jit-able or vmap-able. Matrix exponential is exact for constant generator, JAX-native, and the entire electronic step compiles into XLA.

3.3 The fewest-switches hop probability

Tully’s central result: if you ask “what is the smallest rate of $i \rightarrow j$ hops in the ensemble that keeps the fraction of trajectories on state j equal to $\langle |c_j|^2 \rangle$?”, you get

$$P_{i \rightarrow j}(t \rightarrow t + \Delta t) = \max \left\{ 0, \frac{2 \Delta t \operatorname{Re} [\rho_{ij}^* (\mathbf{v} \cdot \mathbf{d}_{ij})]}{\rho_{ii}} \right\},$$

where $\rho_{ij} = c_i c_j^*$ is the (one-trajectory) density matrix. The numerator is the flux of population *out* of state i along the i - j coupling channel; dividing by ρ_{ii} converts it to a *conditional* transition probability given we’re currently in i .

In `step()` this reads:

```
b = 2.0 * jnp.real(jnp.conj(rho[state.state, :]) * vdotd)
g = jnp.maximum(0.0, b * dt / (rho_ii + 1e-30))
g = g.at[state.state].set(0.0) # never "hop to self"
```

A subtle sign. The derivation in Tully 1990 has terms with both $\mathbf{d}_{ij}$ and $\mathbf{d}_{ji} = -\mathbf{d}_{ij}$ (antisymmetry of the NAC tensor). Depending on which index you “factor out” of the $\operatorname{Re}[\dots]$ trace, you get either a plus or a minus sign in front. The correct convention, verified by reproducing Tully Model 1 transmission within statistical noise, is $b = +2 \operatorname{Re}(\rho_{ij} \mathbf{v} \cdot \mathbf{d}_{ij})$. An early bug in this package had the wrong sign; trajectories ran but the transmission curve was wrong. The Tully-1 benchmark test (`tests/test_tully1_benchmark.py`) guards against this.

Selecting the target state from the probabilities

For $n_{\text{el}} = 2$ the hop decision is just one comparison. For $n_{\text{el}} \geq 3$ Tully’s original algorithm “draw a uniform u , hop to state j if $u < g_{i \rightarrow j}$ ” is *biased*: with three states and equal probabilities to states 1 and 2, comparing the same u to each one in sequence is not the same as drawing one outcome from a categorical distribution. The unbiased version is the **cumulative-probabilities scheme**:

```
cum = jnp.cumsum(g)
u = jax.random.uniform(key, ())
target = jnp.argmax(cum > u) # first bin where cumulative > u
```

This is what `surfacehop_jax` uses. It’s correct for any number of states and reduces to the single- u -comparison for $n_{\text{el}} = 2$.

3.4 Momentum rescaling on a hop

When the trajectory hops from state i to state j , the total energy must be conserved. The classical kinetic energy must absorb the change in potential energy $\Delta E = E_i - E_j$ (positive when going *down*, so kinetic energy increases). Tully prescribes that the rescaling is applied **along the NAC direction**:

$$\mathbf{v}_{\text{new}} = \mathbf{v} - \gamma \frac{\mathbf{d}_{ij}}{m},$$

with γ found by solving the quadratic from energy conservation:

$$\frac{1}{2} \sum_k \frac{d_{ij,k}^2}{m_k} \gamma^2 - (\mathbf{v} \cdot \mathbf{d}_{ij}) \gamma - \Delta E = 0.$$

This has two real roots when the discriminant is non-negative; Tully’s prescription (refined by Hammes-Schiffer & Tully 1994) is to take the root closer to zero (minimal velocity perturbation).

Frustrated hops

If the discriminant is negative, there is no real γ that conserves energy: the trajectory is trying to hop *up* into a state it doesn’t have enough kinetic energy along the NAC to reach. This is a **frustrated hop**. The trajectory stays on its original state, but Truhlar’s 2002 prescription is to **reverse** the velocity component along the NAC:

$$\mathbf{v}_{\text{new}} = \mathbf{v} - \frac{2(\mathbf{v} \cdot \mathbf{d}_{ij})}{\sum_k d_{ij,k}^2 / m_k} \frac{\mathbf{d}_{ij}}{\mathbf{m}}.$$

This reflects the trajectory back toward the coupling region for another chance at a successful hop on a later step, and gives better detailed balance than simply “do nothing”. It is what `surfacehop_jax` implements.

Implementation detail. The momentum-rescaling routine is invoked *unconditionally* on every step (with a placeholder NAC vector of zero when no hop is attempted), then the caller selects the pre-hop velocity if no hop occurred. This avoids a `lax.cond` that XLA can’t fully optimise. The “double-where” pattern around the discriminant prevents $0/0$ from poisoning the autodiff backward pass.

3.5 The eigenvector phase problem

`jnp.linalg.eigh` is an excellent eigensolver, but it returns eigenvectors with **arbitrary signs**: nothing in the eigen problem distinguishes ψ from $-\psi$. As the trajectory walks across the PES, the sign returned at step $t + \Delta t$ may flip independently for each eigenvector. The NAC vectors $\mathbf{d}_{ij}$ depend on the *relative* phases of ψ_i and ψ_j , so an unflagged sign flip would invert the sign of the NAC between steps and turn the trajectory into nonsense.

`pes.fix_eigenvector_phase(new, old)` returns a length-`nel` vector of ± 1 signs such that each new eigenvector best overlaps with the previous one. The next step’s eigenvectors are multiplied by these signs before the NACs are computed. This makes the NAC tensor continuous along the trajectory.

Two failure modes to keep in mind:

1. At an **exact** conical intersection, the eigenvalues are degenerate and `eigh` returns an arbitrary basis. Phase tracking based on single-step overlap fails. In practice for finite-dimensional model Hamiltonians one never lands on the CoIn exactly; if you find trajectories blowing up on a specific step, the time step is probably too large, or you are very near (but not on) a CoIn and should reduce Δt near coupling regions.
2. For trajectories that explore the coordinate widely, the single-step phase-tracking heuristic can drift over very many steps. The propagator is internally consistent — phase is corrected step-to-step — but if you compare absolute eigenvector orientations across thousand-step trajectories, expect arbitrary sign flips. Observables are sign-invariant.

3.6 Internal consistency

A correctly implemented FSSH satisfies the **internal-consistency** property: the fraction of trajectories on adiabatic state j at time t should equal the ensemble-averaged $|c_j(t)|^2$:

$$\frac{1}{N_{\text{traj}}} \sum_{\alpha=1}^{N_{\text{traj}}} \mathbf{1}[s^{(\alpha)}(t) = j] \stackrel{?}{=} \frac{1}{N_{\text{traj}}} \sum_{\alpha} |c_j^{(\alpha)}(t)|^2.$$

This is the whole point of the fewest-switches construction. If your trajectory fraction is systematically far from the average $|c|^2$, something is wrong:

- **Both curves agree but disagree with the exact-quantum reference.** Expected near conical intersections — this is the classic FSSH over-coherence, not an implementation bug. See [Decoherence corrections](#).
- **The two curves diverge as time goes on.** Probable bug: wrong sign in b , missing/wrong frustrated-hop handling, wrong NAC computation, eigenvector phase not tracked.

The Tully-1 benchmark (`tests/test_tully1_benchmark.py`) checks both internal consistency and the transmission curve against Tully’s 1990 published values; it would catch the major implementation bugs in this category.

3.7 What FSSH gets wrong, and what to do about it

The known shortcomings of FSSH, in rough order of severity:

1. **Over-coherence.** The electronic wavefunction stays a coherent superposition long after the underlying nuclear wavepackets would have decohered. Repaired with the Granucci-Persico EDC ([Chapter 8](#)).
2. **Frustrated-hop direction.** Truhlar’s velocity-reflection is the most common prescription but not the only one; for some systems the “do nothing” or “reverse only momentum component” alternatives give different long-time behaviour.
3. **Wavepacket branching information is lost.** Each trajectory makes a hard choice at each hop; the wavepacket character of an actual nuclear wavefunction is not represented. Decoherence corrections mitigate this but don’t eliminate it.
4. **Spatial spread of the wavepacket is not represented.** A single classical trajectory has no width. The Wigner-sampled ensemble recovers the width of the initial state, but coherent quantum effects (interference between branches) are lost.
5. **The classical limit is enforced too aggressively.** Tunneling and zero-point motion are recovered only via the initial Wigner sample; neither develops dynamically.

For a survey of FSSH variants and improvements, see Subotnik et al., *Annu. Rev. Phys. Chem.* **67**, 387 (2016).

Next

Now that the theory is laid out, see how it translates to working code in [04 — Quickstart](#).

Chapter 4: Quickstart

Three runnable examples, each more involved than the last. They mirror the example scripts in [examples/](#). Run any of them in a fresh Python session.

4.1 A single Tully-1 trajectory

Tully's Model 1 is a one-dimensional, two-state avoided crossing. It is the simplest non-trivial FSSH problem.

```
import jax
import jax.numpy as jnp
import surfacehop_jax as sh

model = sh.TullyModel1()
H = model.hamiltonian()          # callable: H(x) -> (2, 2) Hermitian

# Initial conditions: incoming wavepacket on the lower adiabat, k = 15
init = sh.initialize(H,
                    x0=jnp.array([-10.0]),
                    v0=jnp.array([15.0 / 2000.0]), # k / m
                    initial_state=0,
                    nel=2)

final, history = sh.simulate(H, model.masses,
                           init,
                           dt=2.0, n_steps=5000,
                           key=jax.random.PRNGKey(0))

print(f"Final state: {int(final.state)}")
print(f"Final |c|^2 populations: {jnp.abs(final.coeffs)**2}")
print(f"Energy conservation: ΔE = {float(history.total_energy[-1] - history.total_energy[0]):+.}
```

What you should see: the final state is 0 (still on the lower adiabat — transmission) or 1 (jumped to the upper — reflection) depending on the random key, and energy is conserved to better than 10^{-5} Ha over the whole 10 000-step trajectory.

history is a StepDiagnostics named-tuple whose fields are stacked along a leading time axis of length n_steps:

- history.total_energy shape (n_steps,) — useful for energy-conservation tests
- history.population shape (n_steps, nel) — $|c_i|^2$ as a function of time
- history.active_state shape (n_steps,) — the adiabatic state at each step
- history.hopped, history.frustrated — booleans flagging successful and frustrated hop attempts

4.2 A pyrazine ensemble

For a real photochemistry calculation we need (a) a multi-dimensional model, and (b) an ensemble of Wigner-sampled initial conditions. Both are one-liners.

```
import jax
import jax.numpy as jnp
import numpy as np
import surfacehop_jax as sh

model = sh.pyrazine_4mode()      # 2 states × 4 modes
H = model.hamiltonian()

# Wigner-sample 500 initial conditions at the Franck-Condon point
n_traj = 500
```

```

key_qp, key_dyn = jax.random.split(jax.random.PRNGKey(0))
Q0, P0 = sh.sample_phase_space(
    key_qp,
    q0=jnp.zeros(4),          # Franck-Condon = origin
    p0=jnp.zeros(4),
    omega=model.frequencies,
    mass=model.masses,
    n_samples=n_traj,
)
V0 = P0 / model.masses        # velocities

# Build the initial TrajectoryState batch; vmap over trajectories
init = jax.vmap(
    lambda q, v: sh.initialize(H, q, v, initial_state=1, nel=2)
)(Q0, V0)

# Propagate for 120 fs. 1 a.u. of time = 0.0242 fs ⇒ 4960 steps × dt=1 ≈ 120 fs.
final, hist = sh.run_ensemble(H, model.masses, init,
                             dt=1.0, n_steps=4960, key=key_dyn)

# S2 population: trajectory-fraction vs ensemble-averaged |c_S2|²
t = np.arange(4960) * sh.constants.AU_OF_TIME_TO_FS
frac_s2 = (np.asarray(hist.active_state) == 1).mean(axis=0)
pop_s2 = np.asarray(hist.population[:, :, 1]).mean(axis=0)

print(f"P(S2, 60 fs) [active fraction] = {frac_s2[2480]:.3f}")
print(f"P(S2, 60 fs) [mean |c|²] = {pop_s2[2480]:.3f}")

```

The two numbers will not be the same — that’s the well-known FSSH over-coherence near the conical intersection. Turn on a decoherence correction with one extra keyword:

```

from surfacehop_jax.decoherence import zhu_truhlar

final, hist = sh.run_ensemble(H, model.masses, init,
                             dt=1.0, n_steps=4960, key=key_dyn,
                             decoherence_fn=zhu_truhlar)

```

The trajectory-fraction and mean- $|c|^2$ curves should now lie on top of each other (internal consistency restored). See [Chapter 8 — Decoherence](#) for what this does and why.

4.3 A differentiable parameter fit

The headline use case: given an experimental (or hypothetical) target $P(S_2, t = 12 \text{ fs}) = 0.50$, what value of the interstate coupling λ_{10a} reproduces it?

```

import jax
import jax.numpy as jnp
import numpy as np
import optax
import surfacehop_jax as sh
from surfacehop_jax.models import LinearVibronicCoupling

EV = 1.0 / sh.constants.HARTREE_TO_EV
FS_PER_AU = sh.constants.AU_OF_TIME_TO_FS

```

```

def build_pyrazine(lam_eV):
    """Pyrazine 4-mode LVC with  $\lambda_{10a}$  exposed as a JAX-traceable variable."""
    omega = jnp.array([0.07395, 0.12605, 0.15244, 0.09347]) * EV
    energies = jnp.array([3.94, 4.84]) * EV
    kappa_S1 = jnp.array([-0.04634, -0.05382, +0.00795, 0.0]) * EV
    kappa_S2 = jnp.array([+0.10464, +0.04204, +0.05480, 0.0]) * EV
    coupling = jnp.zeros((2, 2, 4))
    coupling = coupling.at[0, 0].set(kappa_S1)
    coupling = coupling.at[1, 1].set(kappa_S2)
    coupling = coupling.at[0, 1, 3].set(lam_eV * EV)
    coupling = coupling.at[1, 0, 3].set(lam_eV * EV)
    return LinearVibronicCoupling(energies=energies, frequencies=omega,
                                   coupling=coupling)

N_TRAJ, T_OBS_FS = 30, 12.0
N_STEPS = int(T_OBS_FS / FS_PER_AU)
KEY = jax.random.PRNGKey(0)

def pS2_at_observe(lam_eV):
    m = build_pyrazine(lam_eV)
    H = m.hamiltonian()
    kq, kd = jax.random.split(KEY)
    Q0, P0 = sh.sample_phase_space(kq, jnp.zeros(4), jnp.zeros(4),
                                   m.frequencies, m.masses, N_TRAJ)

    V0 = P0 / m.masses
    init = jax.vmap(lambda q, v: sh.initialize(H, q, v, 1, 2))(Q0, V0)
    final, _ = sh.run_ensemble(H, m.masses, init, dt=1.0, n_steps=N_STEPS,
                               key=kd, decoherence_fn=sh.decoherence.zhu_truhlar)
    return jnp.mean(jnp.abs(final.coeffs[:, 1]) ** 2)

target = 0.50
@jax.jit
def loss_and_grad(lam):
    return jax.value_and_grad(lambda l: (pS2_at_observe(l) - target) ** 2)(lam)

optimizer = optax.adam(learning_rate=0.05)
lam = jnp.array(0.18)
opt_state = optimizer.init(lam)

for i in range(20):
    loss, g = loss_and_grad(lam)
    updates, opt_state = optimizer.update(g, opt_state, lam)
    lam = optax.apply_updates(lam, updates)
    print(f"iter {i:>2}:  $\lambda = \{{float(lam):.4f}\}$ , loss =  $\{{float(loss):.3e}\}$ ")

```

This whole pipeline — Wigner sample, vmapped ensemble, FSSH propagation with EDC, observable extraction, Adam update — runs end-to-end through `jax.grad`. The same machinery scales to many parameters at no extra forward-pass cost. For the full pedagogical walkthrough with finite-difference validation and convergence plots, open [notebooks/differentiable_dynamics.ipynb](#).

Next

Now that you have something running, the next chapters dig into the pieces:

- [05 — Models](#): the built-in models and how to write your own Hamiltonian.
- [06 — API reference](#): the public functions and what their arguments mean.
- [07 — Wigner sampling](#): correct initial conditions in mass-weighted coordinates.
- [08 — Decoherence](#): when and why to use EDC.
- [09 — Differentiable workflows](#): the full toolkit for `jax.grad`-driven optimisation.

Chapter 5: Models — built-in Hamiltonians and how to write your own

A `surfacehop_jax model` is anything that gives you a callable $H(x) \rightarrow (nel, nel)$ Hermitian matrix. The package ships a handful of canonical ones; rolling your own is one function.

5.1 The model interface

```
@dataclass(frozen=True)
class Model:
    nel: int          # number of electronic states
    ndim: int         # number of nuclear coordinates
    masses: jnp.ndarray # shape (ndim,), in electron-mass units
    name: str = "Model"

    def hamiltonian(self) -> Callable[[jnp.ndarray], jnp.ndarray]:
        ...           # returns a function (ndim,) -> (nel, nel)
```

Every built-in model is a frozen dataclass subclass of `Model` whose `hamiltonian()` method returns a pure JAX function. The propagator doesn't care whether you use the `Model` machinery or pass a raw function — `simulate(H_fn, masses, init, ...)` works either way. The class wrapper is there to keep the parameter set tidy and to make the masses canonical.

5.2 Tully 1990 models

The three canonical 1D benchmarks from Tully's paper. Their parameters match the 1990 publication defaults; the masses are all 2000 (a hydrogen-ish nuclear mass in atomic units).

TullyModel1 — single avoided crossing

$$V_{11}(x) = \text{sgn}(x) A (1 - e^{-B|x|}), \quad V_{22}(x) = -V_{11}(x), \quad V_{12}(x) = C e^{-Dx^2}.$$

Defaults: $A=0.01$, $B=1.6$, $C=0.005$, $D=1.0$. A wavepacket coming in from $x = -\infty$ on the lower adiabat passes through the avoided crossing near $x = 0$; depending on momentum, it either stays on the lower adiabat (low momentum, mostly adiabatic) or jumps to the upper one (high momentum, mostly diabatic). The transmission probability as a function of initial momentum is the classic Tully-1 curve, reproduced by the `tests/test_tully1_benchmark.py` slow test.

```
import surfacehop_jax as sh
model = sh.TullyModel1()           # defaults
model = sh.TullyModel1(A=0.02, C=0.008) # custom A and C
```


Smoothness at $x = 0$. Tully's V_{11} is defined as $\text{sgn}(x) A (1 - e^{-B|x|})$, which is C^1 at the origin (both branches have slope AB) but `jnp.sign(0) = 0` annihilates the autodiff derivative there. The implementation glues two smooth branches with `jnp.where(x >= 0, ..., ...)` so derivatives at the origin are correct.

TullyModel2 — dual avoided crossing

$$V_{11} = 0, \quad V_{22}(x) = -A e^{-Bx^2} + E_0, \quad V_{12}(x) = C e^{-Dx^2}.$$

Defaults: $A=0.10$, $B=0.28$, $E_0=0.05$, $C=0.015$, $D=0.06$. Two avoided crossings on the way through the coupling region produce Stueckelberg oscillations in the transmission curve.

TullyModel3 — extended coupling with reflection

$$V_{11} = A, \quad V_{22} = -A, \quad V_{12}(x) = \begin{cases} B(2 - e^{-Cx}) & x \geq 0 \\ B e^{Cx} & x < 0 \end{cases}.$$

Defaults: $A=6e-4$, $B=0.10$, $C=0.90$. The coupling extends infinitely on the $x > 0$ side, which leads to reflection effects and a richer transmission landscape.

5.3 Linear vibronic coupling

The workhorse model for real photochemistry. The [LinearVibronicCoupling](#) class implements the standard Köppel-Domcke-Cederbaum (KDC) diabatic Hamiltonian as a truncated Taylor expansion around a reference geometry:

$$H_{ij}(\mathbf{Q}) = \delta_{ij} \left[E_i + \frac{1}{2} \sum_{\alpha} \omega_{\alpha} Q_{\alpha}^2 + \sum_{\alpha} \kappa_{\alpha}^{(i)} Q_{\alpha} \right] + (1 - \delta_{ij}) \sum_{\alpha} \lambda_{\alpha}^{(ij)} Q_{\alpha}$$

In words:

- E_i are the vertical excitation energies at the reference geometry (typically a Franck-Condon point).
- ω_{α} are the normal-mode angular frequencies of the ground electronic state.
- $\kappa_{\alpha}^{(i)}$ are the intrastate gradients along mode α on state i — modes with nonzero $\kappa^{(i)}$ are called *tuning modes* because they shift the diabatic minimum of state i .
- $\lambda_{\alpha}^{(ij)} = \lambda_{\alpha}^{(ji)}$ are the interstate couplings — modes with nonzero λ are *coupling modes*.

Coordinates and masses: the dimensionless convention

`LinearVibronicCoupling` uses *dimensionless mass-frequency-weighted normal coordinates*:

$$Q_{\alpha} = \sqrt{\frac{m_{\alpha} \omega_{\alpha}}{\hbar}} q_{\alpha},$$

where q_{α} is the displacement in atomic units of length. In these coordinates the harmonic potential is $\frac{1}{2} \omega_{\alpha} Q_{\alpha}^2$ (notice: ω , not ω^2). The effective classical mass per dimensionless coordinate works out to

$$m_{\text{eff}}^{(\alpha)} = \frac{\hbar}{\omega_{\alpha}},$$

which is what `LinearVibronicCoupling.__post_init__` sets: $\text{masses} = 1 / \text{frequencies}$. **Wigner ground-state widths in these coordinates are $\sigma_Q = \sigma_P = 1/\sqrt{2}$ for every mode**, independent of frequency — a nice consequence of the dimensionless choice that makes Wigner sampling easy.

How to construct one

```
import jax.numpy as jnp
from surfacehop_jax import LinearVibronicCoupling

EV = 1.0 / 27.211386245988          # Hartree per eV

nel, nmodes = 3, 5
energies = jnp.array([0.0, 3.5, 4.2]) * EV      # vertical excitations
omega     = jnp.array([0.08, 0.12, 0.15, 0.18, 0.22]) * EV  # mode frequencies

# Build the (nel, nel, nmodes) coupling tensor. Diagonal: kappa^(i).
# Off-diagonal: lambda^(ij), must be symmetric in (i, j).
coupling = jnp.zeros((nel, nel, nmodes))
coupling = coupling.at[0, 0].set(jnp.array([0.0, 0.0, 0.0, 0.0, 0.0]) * EV)
coupling = coupling.at[1, 1].set(jnp.array([+0.10, +0.03, -0.02, 0.0, 0.0]) * EV)
coupling = coupling.at[2, 2].set(jnp.array([-0.05, +0.06, +0.04, 0.0, 0.0]) * EV)

# Symmetric off-diagonals: lambda_01 along modes 3 and 4
lam_01 = jnp.array([0.0, 0.0, 0.0, 0.08, 0.05]) * EV
coupling = coupling.at[0, 1].set(lam_01)
coupling = coupling.at[1, 0].set(lam_01)

model = LinearVibronicCoupling(
    energies=energies,
    frequencies=omega,
    coupling=coupling,
    name="my-3-state-5-mode-LVC",
)

print(model.nel, model.ndim)          # 3, 5
print(model.masses)                  # 1 / omega per mode

# Get the JAX-traceable Hamiltonian function
H = model.hamiltonian()
print(H(jnp.zeros(5)))                # 3x3 matrix at the FC point
```

Symmetry validation

`LinearVibronicCoupling.__post_init__` checks that the coupling tensor is symmetric in (i, j) — an asymmetric coupling would mean Hermiticity is broken. The check raises `ValueError` if symmetry is violated by more than 10^{-10} .

If you build an LVC **inside a jax.grad- or jit-traced function** (parameter fitting!), the entries of coupling are JAX *tracers*, not concrete arrays, and `float(jnp.max(...))` cannot evaluate them. The symmetry check is wrapped in a `try/except` against `jax.errors.ConcretizationTypeError` so this is a no-op during tracing. **You** are then responsible for passing a symmetric tensor — write your parameterisation to set $[i, j]$ and $[j, i]$ together.

5.4 The pyrazine 4-mode benchmark

`pyrazine_4mode()` returns a fully populated `LinearVibronicCoupling` for the canonical 4-mode pyrazine S_1/S_2 model used throughout the nonadiabatic-dynamics literature:

```
import surfacehop_jax as sh
model = sh.pyrazine_4mode()
print(model.nel, model.ndim)      # 2, 4
print(model)
```

The mode set is:

index	Wilson	ω (eV)	role
0	ν_{6a}	0.0740	tuning
1	ν_1	0.1261	tuning
2	ν_{9a}	0.1524	tuning
3	ν_{10a}	0.0935	coupling (B_{1g})

Vertical excitations: S_1 (1^1B_{3u}) at 3.94 eV, S_2 (1^1B_{2u}) at 4.84 eV. Intrastate gradients (eV per dimensionless Q):

	ν_{6a}	ν_1	ν_{9a}	ν_{10a}
S₁	-0.04634	-0.05382	+0.00795	0
S₂	+0.10464	+0.04204	+0.05480	0

Interstate coupling: $\lambda_{10a} = 0.26152$ eV on the B_{1g} coupling mode only (selection rule: ν_{10a} is the unique mode spanning B_{1g} which is the irrep of $S_2 \otimes S_1 = B_{2u} \otimes B_{3u}$).

The parameter set is the MCTDH-Heidelberg tutorial set, used as the reference for FSSH-vs-MCTDH benchmark comparisons in many papers and the basis of the [pyrazine benchmark figure](#) shipped with the repo.

5.5 Writing a custom Hamiltonian

The model class is convenience scaffolding; the propagator only needs a callable. **Anything that returns a Hermitian (`nel`, `nel`) matrix from a length-`ndim` coordinate is a valid Hamiltonian.** The simplest path is:

```
import jax.numpy as jnp
import surfacehop_jax as sh

def my_H(x):
    # x is shape (ndim,); return shape (nel, nel) Hermitian.
    # JAX-traceable: use jnp, not numpy; use jnp.where for branching.
    e0, e1 = 0.0, 0.1
    coupling = 0.005 * jnp.exp(-x[0]**2)
    return jnp.array([[e0, coupling],
                      [coupling, e1]])

masses = jnp.array([2000.0])

init = sh.initialize(my_H, x0=jnp.array([-5.0]), v0=jnp.array([0.005]),
```

```

        initial_state=0, nel=2)
final, hist = sh.simulate(my_H, masses, init, dt=2.0, n_steps=2000,
                        key=sh.constants.HBAR * 0) # any key works

```

Rules for a custom H

1. **Hermitian.** $H_{ij}(\mathbf{x}) = H_{ji}^*(\mathbf{x})$ for every $\mathbf{x}$. `jnp.linalg.eigh` will return real eigenvalues either way (it symmetrises), but if your H is asymmetric you'll get garbage gradients out of `jacrev` because the adjoint depends on the off-diagonal coupling being a real function of $\mathbf{x}$.
2. **Smooth in $\mathbf{x}$.** No `if x > 0`: Python branching (use `jnp.where`). No NumPy. The propagator backpropagates through H ; any discontinuity breaks autodiff.
3. **Returns a JAX array.** A NumPy array works for the forward pass but breaks `jax.jit`. Use `jnp.array(...)` to build the output.
4. **Pure.** No global mutable state, no I/O, no caches that `jit`-trace will hard-code. Pass anything dynamic in via `jax.tree_util.pytree` closures.

Wrapping in the Model interface

If you want your custom H to live in the Model machinery (so you can call `model.hamiltonian()`, `model.masses`, etc.):

```

from dataclasses import dataclass, field
from typing import Callable
import jax.numpy as jnp
from surfacehop_jax.models import Model

@dataclass(frozen=True)
class MyModel(Model):
    coupling_strength: float = 0.005
    nel: int = 2
    ndim: int = 1
    masses: jnp.ndarray = field(default_factory=lambda: jnp.array([2000.0]))
    name: str = "MyModel"

    def hamiltonian(self):
        C = self.coupling_strength
        def H(x):
            return jnp.array([[0.0, C * jnp.exp(-x[0]**2)],
                              [C * jnp.exp(-x[0]**2), 0.1]])
        return H

```

The `frozen=True` and `field(default_factory=...)` patterns are required by dataclasses for mutable defaults like arrays.

5.6 ML potentials

A natural extension is an ML-fit diabatic Hamiltonian. Because the propagator only requires a JAX-traceable $H(\mathbf{x}) \rightarrow (\text{nel}, \text{nel})$ Hermitian, **any flax/equinox/haiku network whose output you reshape to an (nel, nel) symmetric matrix is a valid Hamiltonian**. The forward trajectory backpropagates straight through the network. This is the seed of a research direction (train an ML diabatic Hamiltonian against MCTDH reference trajectories using gradient descent through FSSH dynamics); the differentiable-fitting notebook is the proof-of-concept.

Next

- [06 — API reference](#) for the public functions.
 - [07 — Wigner sampling](#) for initial conditions in the dimensionless LVC coordinates.
 - [09 — Differentiable workflows](#) for the parameter-fitting recipe.
-

Chapter 6: API reference

Public functions and classes, grouped by module. Every entry includes the call signature, what each argument means, the return type, and a minimal usage example.

6.1 Module structure

```
surfacehop_jax
├── constants      # physical constants
├── models         # diabatic Hamiltonian models
├── pes           # diabatic → adiabatic transform
├── dynamics      # the propagator
├── wigner        # initial-condition sampling
└── decoherence   # decoherence corrections
```

The top-level package re-exports the most-used names:

```
import surfacehop_jax as sh
sh.TullyModel1, sh.TullyModel2, sh.TullyModel3
sh.LinearVibronicCoupling, sh.pyrazine_4mode
sh.TrajectoryState, sh.StepDiagnostics
sh.initialize, sh.step, sh.simulate, sh.run_ensemble
sh.adiabatic_quantities, sh.AdiabaticState
sh.sample_phase_space, sh.wigner_function
sh.constants, sh.decoherence      # submodules
```

6.2 Constants

`surfacehop_jax.constants`:

Name	Value	Notes
HBAR	1.0	$\hbar$ in atomic units
ELECTRON_MASS	1.0	m_e in atomic units
PROTON_MASS	1836.15267343	m_p/m_e , CODATA 2018
HARTREE_TO_EV	27.211386245988	
HARTREE_TO_CM	219474.6313632	
BOHR_TO_ANG	0.529177210903	
AU_OF_TIME_TO_FS	0.024188843265857	

6.3 Models (`surfacehop_jax.models`)

`class Model`

Base class for diabatic Hamiltonian models. Subclass and override `hamiltonian()`.

Attributes:

- `nel`: int — number of electronic states.
- `ndim`: int — number of nuclear coordinates.
- `masses`: (ndim,) array — nuclear masses, electron-mass units.
- `name`: str — human-readable identifier.

Methods:

- `hamiltonian()` -> Callable[[ndim,) array], (nel, nel) array]

Returns a pure JAX function $H(\mathbf{x})$ mapping coordinates to a Hermitian diabatic Hamiltonian.

`class TullyModel1(Model) / TullyModel2 / TullyModel3`

The three Tully-1990 1D benchmarks. See [Chapter 5](#) for definitions. All three default to `masses = jnp.array([2000.0])`.

`class LinearVibronicCoupling(Model)`

Constructor:

```
LinearVibronicCoupling(
    energies: (nel,) array,          # vertical excitations, Hartree
    frequencies: (nmodes,) array,    # mode angular frequencies, Hartree
    coupling: (nel, nel, nmodes) array, #  $\kappa$  on diag,  $\lambda$  off-diag, symmetric in (i,j)
    name: str = "LinearVibronicCoupling",
)
```

Coordinates are *dimensionless mass-frequency-weighted* normal coordinates. Effective per-mode mass $1/\omega$, set automatically.

The `__post_init__` checks shapes and symmetry. The symmetry check is silently skipped when called inside a `jax.grad` or `jax.jit` trace (see [Chapter 9](#)).

`pyrazine_4mode()` -> `LinearVibronicCoupling`

Returns the canonical 4-mode pyrazine S_1/S_2 benchmark with MCTDH-Heidelberg parameter values. No arguments.

6.4 PES (surfacehop_jax.pes)

`class AdiabaticState`

Named tuple of (energies, gradients, nacs, eigvecs) at one point:

- `energies`: (nel,) array — adiabatic eigenvalues, ascending.
- `gradients`: (nel, ndim) array — Hellmann-Feynman gradient $\partial E_i / \partial x_k$.
- `nacs`: (nel, nel, ndim) array — NAC vectors, antisymmetric in (i, j) , zero on the diagonal.
- `eigvecs`: (nel, nel) array — eigenvectors of H , columns indexed by state.

`adiabatic_quantities(diab_h_fn, x)` -> `AdiabaticState`

Diagonalises $H(\mathbf{x})$, computes the gradient tensor via `jax.jacrev`, and assembles energies, gradients, NACs, and eigenvectors in one call. JIT- and vmap-compatible.

```

import jax.numpy as jnp
import surfacehop_jax as sh
H = sh.TullyModel1().hamiltonian()
s = sh.adiabatic_quantities(H, jnp.array([-1.0]))
print(s.energies)           # (2,)
print(s.gradients)          # (2, 1)
print(s.nacs[0, 1])         # NAC vector  $d_{01}$ , shape (1,)

```

6.5 Dynamics (surfacehop_jax.dynamics)

class TrajectoryState

Named tuple of all fields needed to propagate one trajectory by one step:

- `t`: scalar — current time, atomic units.
- `x`: (ndim,) — nuclear positions.
- `v`: (ndim,) — nuclear velocities.
- `state`: scalar int — index of the active adiabatic surface.
- `coeffs`: (nel,) complex — electronic amplitudes (norm 1).
- `energies`, `gradients`, `nacs`, `eigvecs`: cached PES quantities at `x`, identical to the corresponding fields of `AdiabaticState`. Kept on the state so the next step doesn't recompute them.

Being a `NamedTuple`, `TrajectoryState` is a valid JAX pytree without any registration; `jax.jit`, `jax.vmap`, and `jax.lax.scan` all handle it transparently.

class StepDiagnostics

Named tuple of per-step diagnostics, returned alongside the new `TrajectoryState`:

- `hopped`: scalar bool — was there a successful hop this step?
- `frustrated`: scalar bool — was there a frustrated hop attempt?
- `total_energy`: scalar — useful for energy-conservation tests.
- `population`: (nel,) — current $|c_i|^2$.
- `active_state`: scalar int — current adiabatic surface.

When returned from `simulate/run_ensemble`, these fields gain leading time/trajectory axes (see below).

initialize(diab_h_fn, x0, v0, initial_state, nel, *, t0=0.0) -> TrajectoryState

Build the initial state. `coeffs` is set to $\mathbf{e}_{\text{initial_state}}$ (unit population on the initial surface).

```

init = sh.initialize(H,
                    x0=jnp.array([-10.0]), v0=jnp.array([0.0075]),
                    initial_state=0, nel=2, t0=0.0)

```

step(diab_h_fn, masses, state, dt, key, decoherence_fn=None) -> (TrajectoryState, StepDiagnostics)

One velocity-Verlet + TDSE + FSSH step. Pure JAX function — JIT-able, vmap-able. Arguments:

- `diab_h_fn` — callable returning the diabatic Hamiltonian.
- `masses` — (ndim,) nuclear masses.
- `state` — current `TrajectoryState`.

- `dt` — time step, atomic units.
- `key` — `jax.Array` PRNG key, one uniform random number drawn per step for the hop decision.
- `decoherence_fn` — optional, see [surfacehop_jax.decoherence](#). `None` → bare Tully algorithm. Pass `zhu_truhlar` for EDC.

`simulate(diab_h_fn, masses, init_state, dt, n_steps, key, decoherence_fn=None) -> (TrajectoryState, StepDiagnostics)`

Run `n_steps` of FSSH for a single trajectory. Internally uses `jax.lax.scan` so the whole loop compiles to one XLA program. Returns:

- `final_state` — `TrajectoryState` at $t = n \Delta t$.
- `history` — `StepDiagnostics` with **each field's leading axis expanded to length `n_steps`**. So `history.population` is `(n_steps, nel)`, `history.active_state` is `(n_steps,)`, etc.

The history is the standard FSSH output for plotting: average across trajectories at fixed time to get population curves.

`run_ensemble(diab_h_fn, masses, init_states, dt, n_steps, key, decoherence_fn=None) -> (TrajectoryState, StepDiagnostics)`

Run an ensemble of trajectories in parallel via `jax.vmap`. The `init_states` argument is a single `TrajectoryState` whose every field has an additional **leading batch dimension of size `n_traj`**. So if the per-trajectory `init.x` would be `(ndim,)`, the ensemble `init.x` is `(n_traj, ndim)`. The easiest way to construct this is with `jax.vmap(initialize, ...)` over Wigner-sampled `(Q0, V0)` pairs (see [Chapter 7](#)).

The return has the same structure with both batch and time axes:

- `final.x` shape `(n_traj, ndim)`, etc.
- `history.population` shape `(n_traj, n_steps, nel)`.
- `history.active_state` shape `(n_traj, n_steps)`.

To compute ensemble-averaged populations:

```
import numpy as np
pop_S2 = np.asarray(hist.population[:, :, 1]).mean(axis=0)    # (n_steps,)
frac_S2 = (np.asarray(hist.active_state) == 1).mean(axis=0)  # (n_steps,)
```

6.6 Wigner sampling (`surfacehop_jax.wigner`)

`wigner_function(q, p, omega, mass)`

The HO ground-state Wigner distribution, useful for plotting: $W(q, p) = (\pi\hbar)^{-1} \exp(-p^2/m\hbar\omega - m\omega q^2/\hbar)$. Vectorised in q, p .

`sample_phase_space(key, q0, p0, omega, mass, n_samples)`

Draw `n_samples` independent $(\mathbf{q}, \mathbf{p})$ pairs from the multi-mode HO ground-state Wigner distribution. Multi-dimensional: `q0, p0, omega, mass` may all be `(ndim,)` arrays (or scalars). Returns $(\mathbf{q}, \mathbf{p})$ each of shape `(n_samples, ndim)`. See [Chapter 7](#) for the conventions.

6.7 Decoherence (surfacehop_jax.decoherence)

`no_decoherence(coeffs, state, energies, kinetic_energy, dt) -> coeffs`

The identity. Useful as an explicit, named alternative to passing `decoherence_fn=None`.

`zhu_truhlar(coeffs, state, energies, kinetic_energy, dt, alpha=0.1) -> coeffs`

The Zhu-Truhlar form of the Granucci-Persico energy-based decoherence correction. Damps off-active-state amplitudes by $\exp(-\Delta t/\tau_{ij})$ with $\tau_{ij} = (\hbar/|E_i - E_j|)(1 + \alpha/T_{\text{kin}})$, then rescales the active-state amplitude to preserve total norm.

Default $\alpha = 0.1$ Hartree is the value recommended by Truhlar et al. See [Chapter 8](#) for the why, and the original literature ([Granucci & Persico](#), [Zhu et al.](#)) for the derivation.

Plugging a custom decoherence function

The propagator calls `decoherence_fn(coeffs, state, energies, kinetic_energy, dt)` at the end of every step, after the hop decision and momentum rescaling. Any pure JAX function with that signature works:

```
def my_decoherence(coeffs, state, energies, kinetic_energy, dt):  
    # ... return new coeffs (norm-preserving!)  
    return new_coeffs
```

```
final, hist = sh.simulate(H, masses, init, dt, n_steps, key,  
                          decoherence_fn=my_decoherence)
```

Next

- [07 — Wigner sampling](#) for the right way to build ensemble initial conditions.
- [08 — Decoherence](#) for the EDC details.

Chapter 7: Wigner sampling of initial conditions

An ensemble FSSH calculation requires an ensemble of initial conditions that represents a quantum nuclear wavefunction in a classical sense. The standard choice is to draw $(\mathbf{Q}, \mathbf{P})$ pairs from the **Wigner function** of the ground-state nuclear wavefunction. For a multi-mode harmonic oscillator that's a multivariate Gaussian; for anharmonic ground states you'd need a Monte Carlo procedure, but for photochemistry-around-FC the harmonic approximation is standard.

7.1 The ground-state Wigner function

For a one-dimensional harmonic oscillator with mass m and angular frequency ω in the ground vibrational state, the Wigner distribution is the Gaussian

$$W(q, p) = \frac{1}{\pi\hbar} \exp \left[-\frac{p^2}{m\hbar\omega} - \frac{m\omega q^2}{\hbar} \right].$$

Its q - and p -marginals are Gaussian with standard deviations

$$\sigma_q = \sqrt{\frac{\hbar}{2m\omega}}, \quad \sigma_p = \sqrt{\frac{m\hbar\omega}{2}}.$$

The product $\sigma_q \sigma_p = \hbar/2$ saturates the Heisenberg inequality, as expected for the HO ground state.

For multiple uncoupled normal modes (the leading approximation to a real molecule at the FC point) the Wigner distribution factorises and the sampling reduces to drawing independent Gaussians for each mode.

7.2 The API

`sample_phase_space` produces samples from this distribution:

```
import jax
import jax.numpy as jnp
import surfacehop_jax as sh

key = jax.random.PRNGKey(0)
q, p = sh.sample_phase_space(
    key,
    q0=jnp.zeros(4),          # centre of the Gaussian (FC point = 0)
    p0=jnp.zeros(4),          # zero momentum centre
    omega=jnp.array([0.003, 0.005, 0.006, 0.004]), # per-mode frequencies, Hartree
    mass=jnp.array([2000., 2000., 2000., 2000.]),  # per-mode masses, m_e
    n_samples=500,
)
print(q.shape, p.shape)      # (500, 4), (500, 4)
```

Returns positions `q` and momenta `p`. Convert momenta to velocities with $v = p / \text{mass}$ before passing to `initialize`.

7.3 Why this matters: the mass-factor bug

This is the **easiest place to break an FSSH calculation**, and the original PhD code that `surfacehop_jax` rebuilds got this wrong. Walking through it:

The Wigner widths above contain m in *both* the position factor (σ_q) and the momentum factor (σ_p). With $m = 2000$ (roughly hydrogen) and $\omega = 0.01$ Hartree, the correct widths in atomic units are

$$\sigma_q = \sqrt{\hbar/(2 \cdot 2000 \cdot 0.01)} = 0.158, \quad \sigma_p = \sqrt{2000 \cdot 0.01/2} = 3.16.$$

Now suppose you forget the mass in both formulae (the original bug):

$$\sigma_q^{\text{wrong}} = \sqrt{\hbar/2\omega} = 7.07, \quad \sigma_p^{\text{wrong}} = \sqrt{\omega/2} = 0.071.$$

Position too wide by $\sqrt{m} \approx 45$, momentum too narrow by the same factor. The ensemble of initial geometries is hopelessly spread out and the velocities are negligible. Worse, the bug doesn't crash — your trajectories run, they just give wrong populations, slowly. This is the kind of thing a benchmark on a known system catches and that “eyeballing the output” never does.

The repaired `wigner.sample_phase_space` includes the mass factors correctly. The `tests/test_wigner.py` test asserts that the empirical (σ_q, σ_p) from a large sample matches the analytic values to within statistical noise; running it after any future change to the sampler will catch a regression of this bug.

7.4 Sampling in dimensionless coordinates

For a LinearVibronicCoupling model, the package uses *dimensionless mass-frequency-weighted normal coordinates* $Q_\alpha = \sqrt{m_\alpha \omega_\alpha / \hbar} q_\alpha$ and the effective mass is $m_{\text{eff}} = \hbar / \omega_\alpha$. Plugging into $\sigma_q = \sqrt{\hbar / (2m_{\text{eff}}\omega_\alpha)}$:

$$\sigma_Q = \sqrt{\frac{\hbar}{2 \cdot (\hbar/\omega) \cdot \omega}} = \frac{1}{\sqrt{2}}, \quad \sigma_P = \frac{1}{\sqrt{2}}.$$

Both widths are $1/\sqrt{2}$ for every mode, independent of the frequency. This is a nice consequence of the dimensionless choice — the ground-state Wigner blob is an isotropic Gaussian in (Q, P) -space.

When you call `sample_phase_space(key, q0, p0, m.frequencies, m.masses, ...)` on a LinearVibronicCoupling's `(frequencies, masses)`, you get samples in the *dimensionless* coordinates the model lives in. There is no extra conversion to do — just hand them straight to initialize.

7.5 Wigner sampling for an LVC: full pattern

The idiomatic pattern for building an ensemble of TrajectoryStates on an LVC model:

```
import jax
import jax.numpy as jnp
import surfacehop_jax as sh

model = sh.pyrazine_4mode()
H = model.hamiltonian()

n_traj = 500
key_qp, key_dyn = jax.random.split(jax.random.PRNGKey(0))

# 1. Sample positions and momenta in the dimensionless LVC coordinates
Q0, P0 = sh.sample_phase_space(
    key_qp,
    q0=jnp.zeros(model.ndim),          # FC = origin in dimensionless coords
    p0=jnp.zeros(model.ndim),
    omega=model.frequencies,
    mass=model.masses,
    n_samples=n_traj,
)

# 2. Convert momenta to velocities
V0 = P0 / model.masses                 # (n_traj, ndim)

# 3. vmap initialize over the batch. initial_state=1 means we start on S2.
init = jax.vmap(
    lambda q, v: sh.initialize(H, q, v, initial_state=1, nel=2)
)(Q0, V0)

# 4. Now `init` is a TrajectoryState with leading axis n_traj on every field.
final, hist = sh.run_ensemble(H, model.masses, init,
                              dt=1.0, n_steps=4960, key=key_dyn)
```

7.6 Sampling around a non-equilibrium geometry

q_0 and p_0 are the *centres* of the Gaussian distribution. If you want to sample around a displaced geometry — say, a wavepacket prepared at a non-zero coordinate by a chirped pump pulse — just set q_0 to the displaced geometry (in the same coordinate system as the model). The mode widths remain those of the harmonic ground state, which is the right thing if the wavepacket is generated by Franck–Condon excitation of a harmonic ground state.

For a wavepacket of nontrivial width (squeezed states, finite-width pump pulses), or for sampling at thermal equilibrium ($T > 0$), the HO ground-state Wigner is not the right distribution. Build your own sampler that returns (Q, P) of the appropriate shape and feed it into the same `jax.vmap(initialize)` pattern.

7.7 Reproducibility

Wigner sampling uses one PRNG key; the dynamics uses another. **Always split the key explicitly** so the sample and the dynamics realisation can be varied independently. The pattern is `key_qp, key_dyn = jax.random.split(jax.random.PRNGKey(seed))`.

For deterministic gradients in parameter-fitting workflows it’s also useful to *fix* both keys: the gradient is then a well-defined pathwise derivative of the observable with respect to the parameter. See [Chapter 9](#).

Next

- [08 — Decoherence](#) — repair FSSH over-coherence with a one-keyword change.
- [09 — Differentiable workflows](#) — drive Wigner sample + dynamics with `jax.grad`.

Chapter 8: Decoherence corrections

This chapter explains what FSSH over-coherence is, why it matters, what the Granucci-Persico energy-based decoherence correction (EDC) does about it, and how to switch the correction on in `surfacehop_jax`.

If you only want the API: pass `decoherence_fn=surfacehop_jax.decoherence.zhu_truhlar` to `step`, `simulate`, or `run_ensemble`.

8.1 What “over-coherence” means in FSSH

A single FSSH trajectory carries an electronic wavefunction $\Psi(t) = \sum_i c_i(t) \psi_i(\mathbf{Q}(t))$ alongside its nuclear coordinates. After the trajectory passes through a region of strong coupling and the electronic coefficients pick up amplitude on multiple adiabatic states, **the coefficients stay in a coherent superposition indefinitely**. This is just what the TDSE prescribes: in the absence of dissipation, electronic coherence is preserved.

In reality, of course, the *full* nuclear–electronic wavefunction decoheres rapidly. The reason is that the nuclear wavepackets evolving on each electronic surface experience *different* forces — they separate in phase space — and once the nuclear wavepackets on the two surfaces no longer overlap, the off-diagonal elements of the *reduced electronic density matrix* go to zero. The wavefunction has split.

FSSH does not represent this: each trajectory is a single point in nuclear phase space, and the electronic wavefunction propagated along it has no way of “knowing” that the population on the inactive state should be associated with a nuclear wavepacket that is now somewhere

else. The result is **internal inconsistency**: the trajectory fraction on each state ($\langle \mathbf{1}[s = i] \rangle$) and the ensemble-averaged squared amplitude ($\langle |c_i|^2 \rangle$) drift apart. For Tully Models 1-3 the gap is mild. For pyrazine S_2/S_1 near the conical intersection it is dramatic:




Figure 1: Pyrazine bare vs EDC

The left panel (bare FSSH) shows the active-state fraction reaching ~ 0.55 while $\langle |c|^2 \rangle$ on S_2 sits near 0.35 at 120 fs — a 20-percentage-point gap. The right panel (EDC on) shows the two curves overlapping, with the trajectory population on S_2 holding closer to 0.72 at the same time. Internal consistency is restored, and the S_2 lifetime matches high-level quantum reference dynamics on this model.

8.2 The Granucci-Persico EDC: physics

The fix proposed by [Granucci & Persico \(2007\)](#) is conceptually simple: damp the off-active-state amplitudes by an exponential with a physically motivated timescale.

The damping rate is the inverse of the *decoherence time*, which is related to how fast nuclear wavepackets on different electronic surfaces separate. The faster the nuclei (more kinetic energy), the faster the separation, the faster the decoherence. The further apart the surfaces (larger $|E_i - E_j|$), the steeper the force difference, the faster the separation again.

The Granucci-Persico formula combines both effects:

$$\tau_{ij} = \frac{\hbar}{|E_i - E_j|} \left(1 + \frac{C}{T_{\text{kin}}} \right),$$

where T_{kin} is the nuclear kinetic energy and C is a small, fitted “kinetic-energy floor” parameter that keeps τ_{ij} finite when T_{kin} is small. Truhlar and collaborators ([Zhu et al. 2004](#)) found that $C = 0.1$ Hartree works robustly across model and ab initio benchmarks — this is the default in `surfacehop_jax`.

The procedure each step:

1. For every non-active state j , multiply $c_j \rightarrow c_j \cdot \exp(-\Delta t / \tau_{ij})$ where i is the active state.
2. Rescale the active-state amplitude c_i to preserve the total norm $\sum_k |c_k|^2$.

Step 1 reduces the inactive populations toward zero — i.e., toward the classical limit where only the active state is “real”. Step 2 puts the population we just removed back onto the active

state, keeping the wavefunction normalised. The combined effect on a two-state system that started in a pure state is equivalent to multiplying the inactive population by $\exp(-2\Delta t/\tau_{ij})$.

8.3 Using EDC

The decoherence functions live in `surfacehop_jax.decoherence`. Any of them is a drop-in for the `decoherence_fn` argument on `step`, `simulate`, and `run_ensemble`.

```
import surfacehop_jax as sh
from surfacehop_jax.decoherence import zhu_truhlar

# ... build init, H, masses as usual ...

# Bare Tully (default – equivalent to decoherence_fn=None)
final, hist_bare = sh.run_ensemble(H, masses, init, dt=1.0, n_steps=4960,
                                   key=key_dyn)

# Granucci–Persico / Zhu–Truhlar EDC
final, hist_edc = sh.run_ensemble(H, masses, init, dt=1.0, n_steps=4960,
                                   key=key_dyn, decoherence_fn=zhu_truhlar)
```

That’s the whole API.

Plotting internal consistency

The internal-consistency diagnostic uses `hist.population` and `hist.active_state` from the ensemble return. Trajectory fraction is `(active_state == i).mean(axis=0)`; mean $|c_i|^2$ is `population[..., i].mean(axis=0)`. The example script [examples/pyrazine_decoherence.py](#) generates the two-panel comparison above:

```
import numpy as np
import matplotlib.pyplot as plt

t = np.arange(n_steps) * sh.constants.AU_OF_TIME_TO_FS

fig, axes = plt.subplots(1, 2, figsize=(11, 4), sharey=True)
for ax, hist, label in [(axes[0], hist_bare, "bare FSSH"),
                        (axes[1], hist_edc, "with EDC")]:
    frac_s2 = (np.asarray(hist.active_state) == 1).mean(axis=0)
    pop_s2 = np.asarray(hist.population[..., 1]).mean(axis=0)
    ax.plot(t, frac_s2, label="trajectory fraction on S2", lw=2)
    ax.plot(t, pop_s2, label=r"⟨|c_{S2}|2⟩", lw=2, ls="--")
    ax.set_xlabel("time (fs)")
    ax.set_title(label)
    ax.legend()
axes[0].set_ylabel("S2 population")
```

A correctly EDC-corrected ensemble produces two curves that lie on top of each other; a bare ensemble produces two curves that don’t.

8.4 Tuning α (the kinetic-energy floor)

The Zhu–Truhlar formula has one free parameter, α (called C above), with default 0.1 Hartree:

```

from functools import partial
from surfacehop_jax.decoherence import zhu_truhlar

# Try alpha = 0.05 Hartree instead of the default 0.1
my_edc = partial(zhu_truhlar, alpha=0.05)

final, hist = sh.run_ensemble(H, masses, init, dt, n_steps, key,
                             decoherence_fn=my_edc)

```

For most photochemistry problems $\alpha = 0.1$ Hartree is fine — it's the value recommended by the Truhlar group as a one-size-fits-all default. Sensitivity to α is mild as long as it's of order ~ 0.05 – 0.2 Hartree; much smaller and the correction is too strong (kills coherence prematurely, suppresses real recurrences), much larger and it's too weak (over-coherence partially survives).

If you're computing the **sensitivity** of an observable to α , that's just `jax.grad(...)(alpha)` — see [Chapter 9](#).

8.5 When EDC is and isn't necessary

Use EDC for:

- Anything near a conical intersection (pyrazine, cyclohexadiene ring-opening, retinal photoisomerisation, ...). FSSH without decoherence overcoheres dramatically here.
- Any case where you want to compare $\langle |c|^2 \rangle$ to ensemble-averaged trajectory populations: EDC is *required* for internal consistency to hold near regions of strong coupling.
- Long-time photoproduct yields, where small over-coherence per step accumulates into qualitative errors.

EDC may be unnecessary for:

- Tully Models 1–3 with single passage through the avoided crossing. Bare FSSH gets the transmission curve right because the over-coherence is small and there's only one coupling event.
- Calculations where you only need short-time dynamics (the first few fs after photoexcitation) before any decoherence sets in.
- Systems with weak interstate coupling and no CoIn nearby; the bare algorithm is essentially correct.

When in doubt, run with EDC. The cost is a single `jnp.exp` per step; it's noise compared to everything else in the propagator.

8.6 Plug-in custom decoherence

The decoherence interface is the function signature `(coeffs, state, energies, kinetic_energy, dt) -> new_coeffs`. Any norm-preserving JAX function with that signature is a valid correction:

```

import jax.numpy as jnp

def my_decoherence(coeffs, state, energies, kinetic_energy, dt):
    """Coherent switching with decay of mixing (CSDM), simplified."""
    # ... your favourite scheme ...
    return new_coeffs

```

```
final, hist = sh.run_ensemble(H, masses, init, dt, n_steps, key,
                             decoherence_fn=my_decoherence)
```

Implementations to consider if you want to extend the package:

- **Original Granucci-Persico** (with the inactive-state-projection factor $\sum_{j \neq s} |c_j|^2$ in the rescaling). Slightly different from the Zhu-Truhlar form for $n_{\text{el}} \geq 3$.
- **Augmented FSSH** (AFSSH, Subotnik). Carries auxiliary classical-trajectory positions on each surface and collapses the electronic wavefunction when they separate by more than a threshold. Substantially more elaborate; would require an extension to the `TrajectoryState`.
- **DISH** (Decoherence-Induced Surface Hopping, Akimov). A different philosophy: hop *because* of decoherence rather than damping coherence to enforce internal consistency.

For the scope of `surfacehop_jax` (a model-Hamiltonian, differentiable FSSH code), EDC is the right balance of physical realism and implementation simplicity.

Next

- [09 — Differentiable workflows](#) for fitting LVC parameters under EDC.

Chapter 9: Differentiable workflows

This chapter is the long-form companion to the notebook [notebooks/differentiable_dynamics.ipynb](#). It explains *what* “differentiable surface hopping” means, *why* it works at all (and where it doesn’t), and *how* to write your own parameter-fitting workflow on top of `surfacehop_jax`.

9.1 The differentiable propagator: what flows through `jax.grad`

Every line of `surfacehop_jax.dynamics.step` is a pure JAX function. Concretely, the propagator threads a parameter θ (a coupling strength, a frequency, an energy, the weights of an ML potential — anything that goes into the Hamiltonian) through:

1. **The diabatic Hamiltonian** $H(\mathbf{x}; \theta)$ — by construction smooth in θ .
2. **The eigendecomposition** $H\mathbf{v} = E\mathbf{v}$. JAX has reverse-mode rules for eigh that handle degenerate eigenvalues gracefully (well — *non-degenerate* eigenvalues; an exact `Coln` is still a problem in practice, but we never land on one).
3. **The Hellmann-Feynman gradient** $\partial E_i / \partial \mathbf{x}$ and **NACs** via `jax.jacrev` of H . This is itself a derivative, so we’re taking derivatives of derivatives — JAX handles the composition through higher-order autodiff.
4. **Velocity Verlet** for nuclei. Pure arithmetic, trivially differentiable.
5. **The matrix exponential** `jax.scipy.linalg.expm(G * dt)` of the electronic generator G . JAX implements the exact `expm` gradient via Padé approximants.
6. **The hop probability** $g_{i \rightarrow j}$. Smooth in θ .
7. **The hop decision** — *this is the non-smooth step*. The `jnp.argmax(cum > u)` selection is a step function in θ : for most θ no hop happens, then at some threshold the hop threshold crosses u and one does. See §9.4.
8. **Momentum rescaling** on a hop. The “double-where” pattern in `_rescale_velocity_for_hop` prevents $0/0$ in the autodiff backward pass when the rescaling routine is called on a no-hop step.
9. **Decoherence** (`zhu_truhlar`). Pure JAX, fully differentiable.

Compose all of this with `jax.lax.scan` over time steps and `jax.vmap` over an ensemble of trajectories, and you have one big differentiable function from θ to whatever observable of the final `TrajectoryState` you compute. `jax.grad` rebuilds the backward pass through the whole thing.

9.2 The “pathwise” gradient: smooth observables vs hop counts

There’s a subtlety in (7) above worth understanding before you try gradient descent.

The hop decision in FSSH is **discrete**: at each step a hop either happens or it doesn’t, depending on whether $\text{cumsum}(g_{i \rightarrow j}) > u$ for the current random number u . The **whole-ensemble probability of hopping** is smooth in θ — θ slightly larger means g slightly larger means more trajectories hop at a given step. But for any single trajectory with a fixed u , the hop is a step function in θ .

What this means for `jax.grad`:

- For **smooth observables** like $\langle |c_i|^2 \rangle$ (the ensemble-averaged squared amplitude on state i), $\langle Q_\alpha \rangle$ (mean position), kinetic energies, expectation values of any smooth operator — `jax.grad` returns the **pathwise gradient**. This is the derivative *conditional on the realised pattern of hops*: how does the observable change if we infinitesimally tweak θ and keep the same hops? For these observables, pathwise = total gradient up to ensemble noise, because the propagator is smooth between hops.
- For **active-state fractions** like $\langle \mathbf{1}_{[s_\alpha(T) = i]} \rangle$ — *don’t*. The gradient is zero almost everywhere (θ doesn’t change the realised hops), with delta-function spikes at the threshold values of θ where a hop is on the boundary. `jax.grad` returns zero. Use a smooth observable instead.

The pyrazine fit in the notebook uses $\langle |c_{S_2}|^2 \rangle$ at a fixed observation time as the loss target — this is smooth, and the gradient matches finite difference to a few parts in 10^5 .

9.3 Worked recipe: fit one parameter

The full pattern, distilled from the notebook:

```
import jax, jax.numpy as jnp, optax
import surfacehop_jax as sh
from surfacehop_jax.models import LinearVibronicCoupling

EV = 1.0 / sh.constants.HARTREE_TO_EV
FS = sh.constants.AU_OF_TIME_TO_FS

# 1. Differentiable model builder. Note: the parameter (here lam_eV)
# enters as a scalar JAX-traceable value; the rest of the LVC is
# constant.
def build(lam_eV):
    omega = jnp.array([0.07395, 0.12605, 0.15244, 0.09347]) * EV
    energies = jnp.array([3.94, 4.84]) * EV
    kappa_S1 = jnp.array([-0.04634, -0.05382, +0.00795, 0.0]) * EV
    kappa_S2 = jnp.array([+0.10464, +0.04204, +0.05480, 0.0]) * EV
    coupling = jnp.zeros((2, 2, 4))
    coupling = coupling.at[0, 0].set(kappa_S1)
    coupling = coupling.at[1, 1].set(kappa_S2)
    coupling = coupling.at[0, 1, 3].set(lam_eV * EV)
    coupling = coupling.at[1, 0, 3].set(lam_eV * EV)
```

```

    return LinearVibronicCoupling(energies=energies, frequencies=omega,
                                   coupling=coupling)

# 2. Smooth observable function: takes the parameter, returns a scalar.
N_TRAJ, T_OBS_FS = 30, 12.0
N_STEPS = int(T_OBS_FS / FS)
KEY = jax.random.PRNGKey(0)

def loss(lam_eV, target=0.50):
    m = build(lam_eV)
    H = m.hamiltonian()
    kq, kd = jax.random.split(KEY)
    Q0, P0 = sh.sample_phase_space(kq, jnp.zeros(4), jnp.zeros(4),
                                   m.frequencies, m.masses, N_TRAJ)

    V0 = P0 / m.masses
    init = jax.vmap(lambda q, v: sh.initialize(H, q, v, 1, 2))(Q0, V0)
    final, _ = sh.run_ensemble(H, m.masses, init, dt=1.0, n_steps=N_STEPS,
                               key=kd, decoherence_fn=sh.decoherence.zhu_truhlar)
    p_s2 = jnp.mean(jnp.abs(final.coeffs[:, 1]) ** 2)
    return (p_s2 - target) ** 2

# 3. Autodiff loss-and-grad
loss_and_grad = jax.jit(jax.value_and_grad(loss))

# 4. Adam loop
opt = optax.adam(learning_rate=0.05)
lam = jnp.array(0.18) # initial guess
state = opt.init(lam)
for i in range(20):
    l, g = loss_and_grad(lam)
    updates, state = opt.update(g, state, lam)
    lam = optax.apply_updates(lam, updates)
    print(f"iter {i:>2}  lam={float(lam):.4f}  loss={float(l):.3e}")

```

Three things to notice:

- **The PRNG key is fixed** (`KEY = jax.random.PRNGKey(0)`). This makes the gradient a well-defined pathwise derivative; if the key varied with θ the gradient would be dominated by noise. For final validation, run the fitted parameter with *new* keys to check ensemble robustness.
- **The decoherence function is on** (`decoherence_fn=zhu_truhlar`). Fitting against a smooth observable is even smoother once you've removed FSSH over-coherence.
- **The whole thing is JIT-compiled** via the explicit `jax.jit(jax.value_and_grad(...))`. Compile time dominates the first call; subsequent calls re-use the cached XLA program.

9.4 Multi-parameter gradients in one backward pass

The autodiff scaling argument: reverse-mode autodiff computes the gradient of a scalar output with respect to **all inputs** at the cost of one extra backward pass. For a function with n scalar inputs, the gradient costs the same as one forward call, regardless of n . This is the same scaling that makes neural-network training tractable.

For LVCs, the implications are dramatic. Fitting a 24-mode pyrazine model has $O(50)$ LVC parameters (energies, frequencies, κ 's, λ 's). With finite-difference gradients, each Adam

step would require 50+ forward ensemble runs. With autodiff, one forward + one backward — total cost $\sim 2\times$ forward, independent of parameter count.

The pattern is to pack all parameters into a pytree input:

```
params0 = {
    'energies': jnp.array([3.94, 4.84]),
    'kappa_S2': jnp.array([+0.10464, +0.04204, +0.05480, 0.0]),
    'lam_10a': jnp.array(0.26152),
}

def loss(params):
    m = build_from_params(params)      # closure over the model construction
    # ... run ensemble, return scalar loss ...
    return ...

val, grads = jax.value_and_grad(loss)(params0)
# grads is a pytree with the same structure as params0;
# grads['kappa_S2'] is shape (4,), grads['lam_10a'] is a scalar, etc.
```

optax's optimizers all consume pytree-shaped parameters and updates, so the optimisation loop is unchanged from the single-parameter case above except for the line that builds the LVC.

9.5 Validation against finite difference

Before trusting a gradient, *always* compare against centred finite difference at a representative point. A 30-trajectory pyrazine ensemble takes ~ 7 s for value + grad on a CPU (after JIT compile); the FD check is just two extra forward calls.

```
g = jax.grad(loss)(lam0)
eps = 0.002
g_fd = (loss(lam0 + eps) - loss(lam0 - eps)) / (2 * eps)
print(f"autodiff: {float(g):+.5e}")
print(f"FD: {float(g_fd):+.5e}")
print(f"rel err: {abs(float(g) - float(g_fd)) / max(abs(g_fd), 1e-12):.2e}")
```

For the pyrazine λ fit, this gives relative error $\approx 3 \times 10^{-5}$. If your autodiff vs FD disagree by more than $\sim 10^{-3}$, the most likely culprits are:

1. A non-smooth observable (active-state fractions; recover by using $|c|^2$).
2. A non-smooth Hamiltonian (e.g. `jnp.sign(x)` instead of the two-branch `jnp.where`; affects derivatives at the kink).
3. A key that wasn't really fixed (e.g. you used a Python-level RNG somewhere that varies across calls).

9.6 Common pitfalls

ConcretizationTypeError when building a LinearVibronicCoupling inside a traced function

The `__post_init__` symmetry check used to call `float(jnp.max(...))` on the coupling tensor, which fails inside a `jax.grad` trace because the tensor entries are abstract tracers. **This is now wrapped in try / except `jax.errors.ConcretizationTypeError`** so the check is a no-op during tracing — you the user are responsible for passing a symmetric coupling tensor. The standard pattern (set both `[i, j]` and `[j, i]` together) is the easy way to satisfy this.

Gradient is zero everywhere

You're probably differentiating an integer-valued observable like the active-state fraction. Use $|c|^2$ instead.

Gradient is enormous / NaN

Either the propagator hit an exact conical intersection (rare with finite step size; reduce Δt near coupling regions), or you've written a non-smooth Hamiltonian. Sanity check: does the forward pass give sensible numbers? Does FD give sensible numbers? If FD is fine but autodiff is enormous, that's a JAX bug or a numerical instability; report it.

Compile time dominates wall time

Every distinct call signature (shapes, dtypes, Python-level constants) to a jit-ed function recompiles. Avoid passing scalars as int or float literals; wrap them in `jnp.array(...)` so they're traced as arrays. The `dt` argument to `step/simulate/run_ensemble` is fine either way (it's a Python float), but if you find unexpected recompiles, this is the first thing to check.

`n_steps` makes the trace explode

`jax.lax.scan(body, init, xs)` traces `body` once and the loop runs n times in XLA, but the **trace cost itself** grows with the depth of any nested transformations (`vmap` of `scan` of `value_and_grad` of `vmap` ...). On CPU with no GPU acceleration, ensembles of ~ 30 trajectories $\times \sim 500$ steps compile in ~ 10 s and run a forward+backward pass in ~ 7 s. Larger problems compile longer; consider profiling with `jax.profiler` if compile time becomes intolerable.

9.7 Beyond LVC fits: ideas

The same machinery generalises beyond fitting LVC parameters to a target observable. Sketches:

- **Inverse design of a photoswitch.** Define the model as an LVC plus a tunable substituent term; let the parameter be the substituent's electronic effect on the diabatic energies; minimise a loss that combines a target S_2/S_1 branching ratio with a synthetic accessibility penalty.
- **Training an ML diabatic potential against quantum reference dynamics.** Replace the LVC builder above with a flax/equinox neural network; treat its weights as the parameters; the loss is the L^2 distance from `surfacehop_jax` ensemble dynamics to a high-level reference (MCTDH / multi-configuration Ehrenfest) at a set of trajectory snapshots. `jax.grad` returns the gradient of the loss through the FSSH propagator and into the network weights.
- **Hamiltonian-aware HMC over LVC posteriors.** Use `jax.grad` of a log-likelihood (FSSH-observable vs experimental TR-PES) plus a prior to drive a Hamiltonian Monte Carlo sampler over the LVC parameter space. Probabilistic LVC fits.

All of these are PhD-thesis-scale projects, not one-off scripts. But the substrate is here: `jax.grad` flows through every line of the propagator, on CPU and GPU, exactly as it does through a JAX neural network.

Next

- [10 — Performance](#) for scaling to large ensembles and long trajectories.
- [11 — Troubleshooting](#) for the longer list of gotchas.

Chapter 10: Performance — JIT, vmap, GPU

The whole reason `surfacehop_jax` is built on JAX is that the propagator becomes one XLA program and an ensemble of trajectories runs as one batched program. This chapter explains how that scales, what dominates wall time, and how to use the package on a GPU.

10.1 The cost model

For an ensemble of N trajectories propagated for T steps with n_{el} electronic states and n_{dim} nuclear coordinates, the wall-clock cost is roughly:

$$\text{wall time} \approx \underbrace{T_{\text{compile}}}_{\text{one-shot}} + T \cdot N \cdot c(n_{\text{el}}, n_{\text{dim}})$$

where $c(\cdot)$ is the per-step cost of the propagator on one trajectory. The compile time is paid once per shape signature; thereafter the XLA program runs at full throughput. The per-step cost scales as $O(n_{\text{el}}^3)$ from the eigendecomposition and $O(n_{\text{el}}^2 n_{\text{dim}})$ from the gradient tensor; for typical photochemistry ($n_{\text{el}} = 2\text{--}5$, $n_{\text{dim}} = 4\text{--}30$) these are tiny by modern standards.

What this means in practice (CPU laptop, 2024 hardware):

problem	N traj $\times$ T steps	wall (forward)	wall (forward+backward)
Tully-1, full transmission curve	500×5000	~ 5 s	n/a
Pyrazine 4-mode, 120 fs	500×5000	~ 30 s	n/a
Pyrazine fit, one gradient eval	30×500	~ 3 s	~ 7 s

GPU performance depends on the device, but transparently runs the same code: ensembles of 5000+ trajectories at 5000 steps each are comfortable on a single consumer GPU.

10.2 Why JIT helps

`jax.lax.scan` traces the step function **once** and the resulting XLA program runs the loop natively. Without `scan`, a Python for loop calling `step` would have one Python round-trip per step — the overhead per step would be ~ 50 μs and the actual propagator work would be a few μs , i.e., 90%+ overhead.

With `scan`, you pay Python overhead exactly once (during tracing) and then the loop is XLA-native. For long trajectories this is the difference between minutes and seconds.

The `simulate` and `run_ensemble` functions both wrap their internal loops in `scan`, so as long as you call them at the function level (not piece-by-piece via `raw_step`), you get this for free.

10.3 Why vmap helps even more

`jax.vmap` is the second multiplier. The pattern for ensembles is:

```
final, hist = sh.run_ensemble(H, masses, init, dt, n_steps, key)
```

which internally is

```
def one_traj(init, k):  
    return simulate(H, masses, init, dt, n_steps, k, decoherence_fn)  
final, hist = jax.vmap(one_traj)(init_states, keys)
```

`vmap` adds a leading axis to every operation in the traced program; the entire ensemble runs as one batched call. On CPU you get vectorisation; on GPU you get hundreds-to-thousands-way parallelism nearly for free.

Important: every field of `init_states` must have a leading `n_traj` axis. The standard Wigner-sample-and-initialise pattern from [Chapter 7](#) produces this naturally:

```
Q0, P0 = sh.sample_phase_space(key_qp, ..., n_samples=n_traj)  
V0 = P0 / masses  
init = jax.vmap(lambda q, v: sh.initialize(H, q, v, 1, 2))(Q0, V0)
```

If you ever need to assemble the ensemble by hand, use `jax.tree.map(lambda x: jnp.stack([x_traj_i for ...], axis=0), template)` to add the batch axis to every field.

10.4 Compile time and trace caching

The first call to a jit-compiled function pays the compile cost. The second call with the **same** argument shapes and dtypes reuses the cached XLA program. The compile cost scales with the depth of any nested transformations and the size of the unrolled trace; for the pyrazine ensemble (30 traj, 500 steps, `vmap` of `scan` of `expm` of `eigh`) it's ~10 seconds on a laptop CPU.

Things that *trigger a recompile*:

- Changing `n_steps`. `lax.scan` is shape-polymorphic in the loop count, but if you change `n_steps` between calls, JAX has to re-trace because the step *count* is part of the type signature.
- Changing `n_traj`. Same reason: the leading `vmap` axis is part of the type signature.
- Switching `decoherence_fn` between `None` and `zhu_truhlar` (or any other change to the Python-level structure of the function).
- Changing `dt` if it's a Python scalar (JAX promotes Python floats to concrete constants in the trace). Wrap as `jnp.array(dt)` to make it traced.
- Changing dtype (e.g., toggling `x64`).

Things that *don't* trigger a recompile:

- Changing the values of `init`, `key`, or any other JAX-array argument while keeping shapes the same.
- Changing the parameters threaded through the Hamiltonian builder (this is the whole point of differentiable workflows).

For parameter-fitting loops, lock in the shapes by establishing them on the first call and reusing them thereafter.

10.5 GPU

Everything in `surfacehop_jax` is `vmap`- and `jit`-compatible, so it runs on whatever device your `jaxlib` is configured for. The way to move computation to the GPU is to make sure your input arrays live on the GPU device when you call `run_ensemble`. JAX does this automatically when the GPU is the default device:

```
import jax
print(jax.devices())
# [CudaDevice(id=0)] ← GPU is default

# Build inputs as usual; they go to the GPU.
final, hist = sh.run_ensemble(H, masses, init, dt=1.0, n_steps=5000,
                             key=jax.random.PRNGKey(0))
# Outputs are on the GPU. hist.population is a GPU array.
```

For multi-GPU runs you'd need `jax.pmap` (or the newer `jax.sharding` API) wrapped around `run_ensemble`; the package doesn't ship this because the single-GPU regime is large enough for most model-Hamiltonian problems. The single-GPU pattern: build an ensemble of $N \cdot n_{\text{shards}}$ trajectories and run one call; the GPU saturates with $N \gtrsim 1000$ for $n_{\text{el}} = 2$ problems and $N \gtrsim 100$ for $n_{\text{el}} \geq 5$.

10.6 Memory

For an ensemble of N trajectories propagated for T steps, the `hist` return contains:

- `hist.population` shape (N, T, n_{el}) , `float64`
- `hist.active_state` shape (N, T) , `int64`
- `hist.hopped`, `hist.frustrated` shape (N, T) , `bool`
- `hist.total_energy` shape (N, T) , `float64`

For a 1000-trajectory $\times$ 5000-step $\times$ 2-state run, the dominant arrays are $1000 \times 5000 \times 2 \times 8$ bytes = 80 MB for population alone. Not huge, but on GPU you can hit memory limits with bigger ensembles. Two mitigations:

1. **Don't return the history.** If you only need the final state, you can write a thin wrapper around `simulate` that discards the history output inside the scan body. This drops the memory cost to $O(N \cdot n_{\text{state_size}})$, independent of T . The package doesn't provide this convenience function because most diagnostic plots need the time history; you can add it locally.
2. **Subsample the history.** Write a scan that returns the diagnostic only every k steps (still write the full propagator, just yield from body conditionally). This is a `lax.scan` with a custom carry that includes a step counter.

For parameter-fitting workflows that compute one scalar loss per forward pass, neither matters: discard `hist`, take just the relevant field of `final`.

10.7 When FSSH gets slow

The package is bottlenecked by:

1. `jnp.linalg.eigh` on the diabatic Hamiltonian. $O(n_{\text{el}}^3)$ per step. For $n_{\text{el}} \lesssim 10$ this is fast; for $n_{\text{el}} \sim 100$ (some condensed-phase models) it would start to dominate.
2. `jax.scipy.linalg.expm` on the electronic generator. Also $O(n_{\text{el}}^3)$ per step, but typically with a smaller constant than `eigh`.

3. **jax.jacrev** of the Hamiltonian. For an LVC this is fast (the Hamiltonian is a single ein-sum + diag). For an ML potential with millions of parameters it's the bottleneck and dominates everything else.

For LVC problems the propagator is essentially free; for ML-PES problems the network evaluation dominates and there's not much the FSSH layer can do about it. In both cases the JIT-compiled scan + vmap structure is what you want.

10.8 Profiling

To see where time is going, wrap a small representative call in `jax.profiler`:

```
import jax
with jax.profiler.trace("./tmp_profile"):
    final, hist = sh.run_ensemble(H, masses, init, dt=1.0, n_steps=500,
                                key=jax.random.PRNGKey(0))
```

The trace can then be loaded in TensorBoard's profiler UI for a flame-graph view. Alternatively, `%timeit` in a Jupyter notebook is often enough — call the function once to compile, then `%timeit` to measure steady-state cost.

Next

- [11 — Troubleshooting](#) for diagnosing specific failure modes.
-

Chapter 11: Troubleshooting

This chapter is a long-form FAQ for things that have actually broken during development of the package, in roughly the order they're likely to surface. If your problem isn't here, the test suite is the second place to look — every public function has at least one test, and the tests are short enough to read in a few minutes.

11.1 Setup and import issues

jnplinalg.eigh complains about complex numbers

You probably have float32 JAX. `surfacehop_jax` enables float64 on import via

```
import jax
jax.config.update("jax_enable_x64", True)
```

in the package's `__init__.py`. If something else in your code (or an import order quirk) disables it after we've enabled it, eigenvalue problems on the electronic generator can underflow / lose precision. Put `import surfacehop_jax` early, ideally before any other JAX-using imports.

ImportError: cannot import name 'jaxlib'

`pip` resolved `jax` to a newer version than `jaxlib` can supply on your platform. Pinning both versions explicitly usually fixes it:

```
pip install "jax==0.4.30" "jaxlib==0.4.30"
```

Then reinstall `surfacehop_jax` on top.

GPU not detected

```
import jax
print(jax.devices())    # [CpuDevice(id=0)] when you wanted CUDA
```

Most common cause: a CUDA-toolkit/driver mismatch. JAX's install page has a compatibility table. Less common: you installed both the CPU and CUDA wheels and the CPU one was picked. `pip uninstall jax jaxlib` and reinstall just the CUDA wheel.

11.2 Forward-pass problems

Energy isn't conserved

A few likely causes:

1. **The time step is too large for the gradient magnitude.** FSSH is energetically exact in the limit $\Delta t \rightarrow 0$; for typical problems $\Delta t \approx 1$ a.u. = 0.024 fs is the upper end of what's safe with the matrix-exponential propagator. If energy drifts noticeably (more than 10^{-4} Hartree over hundreds of steps), halve Δt .
2. **Frustrated hops with velocity reversal.** Energy is exactly conserved across a successful hop (by construction) and across a frustrated hop (the velocity reflection is symplectic), but if you inspect *only* the kinetic-plus-potential of the active state, a frustrated hop looks like a velocity flip with no force update, and the *next* step's velocity-Verlet update will recover the right energy. So look at `history.total_energy[-1] - history.total_energy[0]` over the whole trajectory, not differences across single steps near frustrated hops.
3. **Non-Hermitian Hamiltonian.** If your custom $H(x)$ is asymmetric, `jnp.linalg.eigh` returns real eigenvalues anyway (it implicitly symmetrises), but the energy you compute via the eigenvectors doesn't match what the propagator integrates. Make H Hermitian.

Trajectories sit forever on the initial state

The hop probability g is being computed correctly, but always coming out near zero. The probable cause is the **velocity used in `vdotd`**: if v is essentially zero (e.g., the trajectory started at rest at the FC point in a high-symmetry coordinate system), the $v \cdot d_{ij}$ term in the generator is zero and population never transfers off the initial state, so no hops happen. Either start with non-zero velocity (Wigner sample with $p_0 \neq 0$, or perturb), or verify that the Hamiltonian actually has an inter-state coupling at the initial geometry.

Sudden enormous force / position diverges

The trajectory walked too close to a conical intersection and `eigh` returned ill-defined eigenvectors. Mitigations:

- Reduce Δt globally, especially near coupling regions.
- Check whether your Hamiltonian has a real CoIn (eigenvalue degeneracy) at any geometry the trajectory can reach; if so, the current propagator will have trouble exactly there.
- Inspect the trajectory: `print(x, v, energies)` every few steps near the divergence and identify the step where it goes wrong.

Eigenvector sign flips visible in plots

Plots of the eigenvector components or NAC vectors show sudden sign flips between steps. The propagator's internal phase tracking (`fix_eigenvector_phase`) keeps phases consistent

step-to-step inside the trajectory, but it doesn't fix global orientation. If you plot raw NAC components across a long trajectory you'll see flips where the phase tracking was indifferent (e.g., near an avoided crossing where eigenvector continuity isn't unique).

The trajectory itself is fine — observables and energies are sign-invariant. If you're plotting NAC components for diagnostics, plot $|NAC|$ instead.

11.3 Ensemble problems

Population curves are noisy

Either you don't have enough trajectories ($N \lesssim 100$ for a photochemistry problem gives visibly noisy curves), or you're computing the *active-state fraction* of a poorly-mixed problem (the fraction is inherently noisier than $\langle |c|^2 \rangle$ for small N). Use more trajectories or plot population instead of `active_state`.

Trajectory fraction and mean $|c|^2$ disagree

This is **FSSH over-coherence**, not a bug — it's the standard symptom near a conical intersection. Switch on the Granucci-Persico EDC ([Chapter 8](#)).

If the curves disagree even *with* EDC on, the implementation may have a real bug. Sanity checks:

- Tully Model 1 with bare FSSH: trajectory fraction and $\langle |c|^2 \rangle$ should agree to within 10^{-2} over 5000 steps (the tests/ Tully-1 test verifies this).
- Pyrazine 4-mode with EDC on: trajectory fraction and $\langle |c|^2 \rangle$ should overlap visually for the first ~ 100 fs.

If either fails, report it; we'd want to dig in.

`run_ensemble` fails with shape error

The most common message is something like `TypeError: dot_general requires contracting dimensions to have ... or similar`. The cause is almost always **`init_states` doesn't have a leading `n_traj` axis on every field**. Check:

```
print(init.x.shape)           # should be (n_traj, ndim)
print(init.coeffs.shape)      # should be (n_traj, nel) complex
print(init.state.shape)       # should be (n_traj,)
```

The correct way to build this is `jax.vmap(initialize, in_axes=...)` over the Wigner-sampled ($Q0$, $V0$) arrays; see [Chapter 7](#).

11.4 Autodiff problems

`ConcretizationTypeError` when building `LinearVibronicCoupling`

In versions ≤ 1.0 the LVC `__post_init__` symmetry check called `float(jnp.max(...))` on the coupling tensor, which fails when the tensor is a JAX tracer (i.e., we're inside `jax.grad` or `jax.jit`). **Fixed in v1.1:** the check is wrapped in `try / except jax.errors.ConcretizationTypeError` so it's a no-op during tracing. You're now responsible for passing a symmetric coupling tensor — the standard pattern is to set `coupling[i, j]` and `coupling[j, i]` together.

If you still see this error, you're running an old version (`pip show surfacehop_jax`). Upgrade.

jax.grad returns zero

You're differentiating a step-function observable — almost always the active-state fraction (`active_state == i`).mean(). Active-state fractions involve discrete hop decisions, so their derivative is zero almost everywhere. Use a smooth observable like `jnp.mean(jnp.abs(coeffs[:, i])**2)` instead.

jax.grad returns NaN

Most often: the trajectory hit a singular point (an exact CoIn) or a $0/0$ ratio in the momentum-rescaling routine. The latter is guarded against in the implementation (the “double-where” pattern around the discriminant), so check the former: print the energies at every step and see if any pair is anomalously close.

If the NaN persists, the culprit is sometimes the matrix exponential on a generator with very large $\|G\| \Delta t$. Reduce Δt .

Autodiff and finite difference disagree

Run the sanity check:

```
g = jax.grad(loss)(theta0)
eps = 1e-3
g_fd = (loss(theta0 + eps) - loss(theta0 - eps)) / (2 * eps)
```

If they disagree by more than $\sim 1\%$, common causes are:

1. **Non-smooth observable:** use $|c|^2$, not active-state fraction.
2. **Non-smooth Hamiltonian:** `jnp.sign(x)` instead of two-branch `jnp.where`; `jnp.heaviside`; Python `if` clauses depending on the coordinate.
3. **PRNG key not fixed:** the autodiff gradient is pathwise (one realisation of the random hops); if the FD is computed with a different key, the two answers come from different paths and disagree.

Compile time is intolerable

First call to a function takes minutes. The trace got large — common causes:

- **n_steps or n_traj is very large at first call.** Start with a smaller call to warm up the cache, then scale up.
- **You're differentiating through a vmap of a scan.** This is fine but the trace is deeper. Wrap with `jax.jit` so the trace cost is paid once.
- **The Hamiltonian itself is large.** If your H is a neural network with millions of parameters, `jacrev` will dominate trace time. Consider whether you really need the full Jacobian or could use a forward-mode `jacfwd` instead.

11.5 Numerical / accuracy issues

Tully Model 1 transmission curve has the wrong shape

If the curve is *flipped* (high momentum stays on the lower adiabat, low momentum hops to the upper), the sign of the b term in the hop probability is wrong. The correct formula is $b = +2 \operatorname{Re}(\rho_{ij}^*(v \cdot d_{ij}))$. The tests/test_tully1_benchmark.py test catches this.

Pyrazine S₂ population doesn't match the literature

A few possible reasons:

1. **You're plotting the wrong thing.** "Population on S₂" can mean the active-state fraction or the mean $|c|^2$; the two are equal with EDC and unequal without. The literature canonical curves ([Worth, Meyer, Cederbaum 1998](#) etc.) are exact-quantum wavepacket calculations and don't have this distinction.
2. **You're using the wrong parameter set.** Several "pyrazine 4-mode" parameter sets exist in the literature: Schneider & Domcke 1988, Worth/Meyer/Cederbaum 1998, Raab et al. 1999, the MCTDH-Heidelberg tutorial. `pyrazine_4mode()` uses the MCTDH-Heidelberg tutorial set. If you're comparing to a paper using a different set, expect quantitative differences.
3. **You forgot EDC.** Bare FSSH overcoheres for pyrazine; the active-state fraction will not match the literature curves quoted from exact-quantum calculations.

Test `test_wigner.py` fails after a refactor

The most likely cause is a regression of the mass-factor bug discussed in [Chapter 7](#). The widths (σ_q, σ_p) should each contain a factor of $\sqrt{m}$. The test sanity-checks empirical against analytic widths to within a few percent.

11.6 Getting help

This package is a research tool, and the test suite and documentation are the primary support resources. If you find a real bug, the maintainer welcomes a minimum-working-example issue or PR.

Things that **are** in scope to fix:

- Numerical bugs (wrong signs, wrong rescaling, wrong NAC formula).
- Autodiff regressions (a `ConcretizationTypeError`, a NaN gradient in a place it shouldn't).
- Performance regressions (a recompile that shouldn't happen, a memory blowup).

Things that **aren't** in scope:

- "Why does FSSH disagree with experiment?" That's photochemistry, not software.
- New methods (CSDM, AFSSH, DISH). Pull requests welcome; new modules alongside `decoherence.py` are the right pattern.
- On-the-fly ab initio interfaces. Use SHARC or Newton-X.

About this guide

Generated from the per-chapter markdown files in `docs/`. For the source-of-truth multi-file version (with proper chapter-by-chapter navigation on GitHub), see [docs/](#). For the JOSS publication summary, see [joss/paper.md](#). For the live API, import `surfacehop_jax`; `help(surfacehop_jax)`.